

# **W271-2 – Spring 2016 – Lab 2**

Juanjo Carin, Kevin Davis, Ashley Levato, Minghu Song

March 7, 2016

## **Contents**

<b>Question 1: Broken Rulers</b>	<b>2</b>
<b>Question 2: Investing</b>	<b>8</b>
<b>Question 3: Turtles</b>	<b>10</b>
<b>Question 4: CLM 1</b>	<b>12</b>
Background . . . . .	12
The Data . . . . .	12
Question 4.1 . . . . .	12
Question 4.2 . . . . .	15
Question 4.3 . . . . .	18
Question 4.4 . . . . .	20
Question 4.5 . . . . .	21
Question 4.6 . . . . .	27
<b>Question 5: CLM 2</b>	<b>30</b>
<b>Question 6: CLM 3</b>	<b>41</b>

---

## Question 1: Broken Rulers

You have a ruler of length 1 and you choose a place to break it using a uniform probability distribution. Let random variable  $X$  represent the length of the left piece of the ruler.  $X$  is distributed uniformly in  $[0, 1]$ . You take the left piece of the ruler and once again choose a place to break it using a uniform probability distribution. Let random variable  $Y$  be the length of the left piece from the second break.

- Find the conditional expectation of  $Y$  given  $X$ ,  $E(Y|X)$ .

$f_X = U(0, 1)$  and  $f_{Y|X} = U(0, X)$  (because the maximum length of the second left piece cannot be greater than the length of the first left piece). As we know, the probability density function for a variable  $Z$  that follows a uniform distribution  $U(a, b)$  is:

$$f_Z(z) = \begin{cases} \frac{1}{b-a} & a \leq z \leq b \\ 0 & \text{otherwise} \end{cases}$$

So:

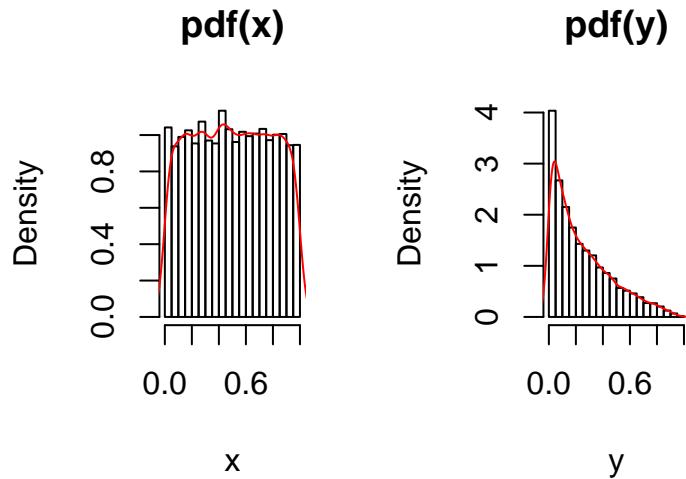
$$f_{Y|X}(y|x) = \begin{cases} \frac{1}{x} & 0 \leq y \leq x \\ 0 & \text{otherwise} \end{cases}$$

And:

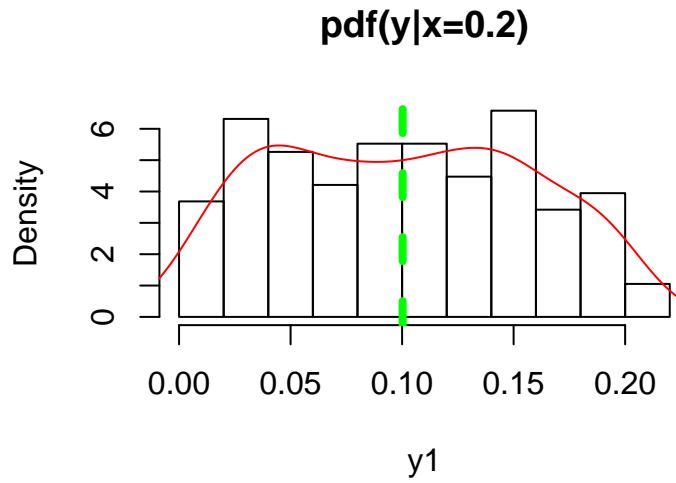
$$\mathbf{E}(Y|X) = \int_{\mathbb{Y}} y \cdot f_{Y|X}(y|x) \cdot dy = \int_{y=0}^x y \cdot \frac{1}{x} \cdot dy = \frac{1}{x} \left[ \frac{y^2}{2} \right]_0^x = \frac{x^2}{2x} = \frac{x}{2}$$

We'll make use of some simulations through this Question to confirm the results.

```
simulations <- 1e4 # number of simulations
set.seed(123)
x <- runif(simulations, min=0, max=1) # X ~ U(0, 1)
y <- runif(simulations, min=0, max=x) # Y/X ~ U(0, X)
par(mfrow = c(1, 2))
hist(x, main = "pdf(x)", freq = FALSE)
lines(density(x), col = 'red')
hist(y, main = "pdf(y)", freq = FALSE)
lines(density(y), col = 'red')
```

Figure 1: Histogram and approximate pdf of  $X$  and  $Y$ 

```
# y1 <- runif(simulations, min = 0, max = 0.2) # Fix X to 0.2
y1 <- y[x > 0.2 - 1e-2 & x < 0.2 + 1e-2] # Using previous simulation
hist(y1, main = 'pdf(y|x=0.2)', freq = FALSE)
lines(density(y1), xlim = c(0, 1), main = 'pdf(y|x=0.2)', col = 'red')
abline(v = mean(y1), col = 'green', lty = 2, lwd = 4)
```

Figure 2: Histogram and approximate pdf of  $Y$  conditional on  $X$  for a given value of  $X$  (0.2)

```
# legend("topright", "E(Y|X=0.2)", lty = 1, bty="n", col = 'red')
```

2. Find the unconditional expectation of  $Y$ . One way to do this is to apply the law of iterated expectations, which states that  $E(Y) = E(E(Y|X))$ . The inner expectation is the conditional expectation computed above, which is a function of  $X$ . The outer expectation finds the expected value of this function.

$$\mathbf{E}(\mathbf{Y}) = E[E(Y|X)] = \int_{\mathbb{X}} E(Y|X) \cdot f_X(x) \cdot dx = \int_{x=0}^1 \frac{x}{2} \cdot 1 \cdot dx = \left[ \frac{x^2}{4} \right]_{x=0}^1 = \frac{1}{4} = \mathbf{0.25}$$

3. Write down an expression for the joint probability density function of  $X$  and  $Y$ ,  $f_{X,Y}(x,y)$ .

$$f_{X,Y}(x,y) = f_{Y|X}(y|x) \cdot f_X(x) = \begin{cases} \frac{1}{x} & x \in (0,1), y \in (0,x) \\ 0 & \text{otherwise} \end{cases}$$

Let's check that this is a valid joint *pdf*:

$$\int_{\mathbb{X}} \int_{\mathbb{Y}} f_{X,Y}(x,y) \cdot dx \cdot dy = \int_{x=0}^1 \int_{y=0}^x \frac{1}{x} \cdot dy \cdot dx = \int_{x=0}^1 \left[ \frac{y}{x} \right]_{y=0}^x dx = \int_{x=0}^1 dx = [x]_{x=0}^1 = 1$$

Simulations:

```

pdf_x <- function(x) ifelse(x<1 & x>0, 1, 0) # f(x)
integrate(pdf_x, -Inf, Inf) # integral

## 1 with absolute error < 4.2e-11

pdf_y_given_x <- function(x,y) ifelse(y<x & y>0 & x<1 & x>0, 1/x, 0) # f(y/x)
pdf_xy <- function(x,y) pdf_x(x)*pdf_y_given_x(x,y) # f(x,y)
# integral
integrate(function(y) sapply(y, function(y) integrate(function(x)
  pdf_xy(x,y), 0, 1)$value), -Inf, Inf)

## 1 with absolute error < 9.3e-05

# Plot f(x,y)
x0 <- y0 <- seq(0, 1, by = 0.01)
grid <- mesh(x0, y0)
z0 <- with(grid, pdf_x(x)*pdf_y_given_x(x,y))
# contour(x0, y0, z0, asp=1)
# par(mfrow = c(1, 2))
persp3D(z = z0, x = x0, y = y0)

# Confirm that f(x,y) = 1/x
# z2 <- with(grid, ifelse(x<=y | x==0 | y == 0, 0, 1/x))
# persp3D(z = z2, x = x0, y = y0)

```

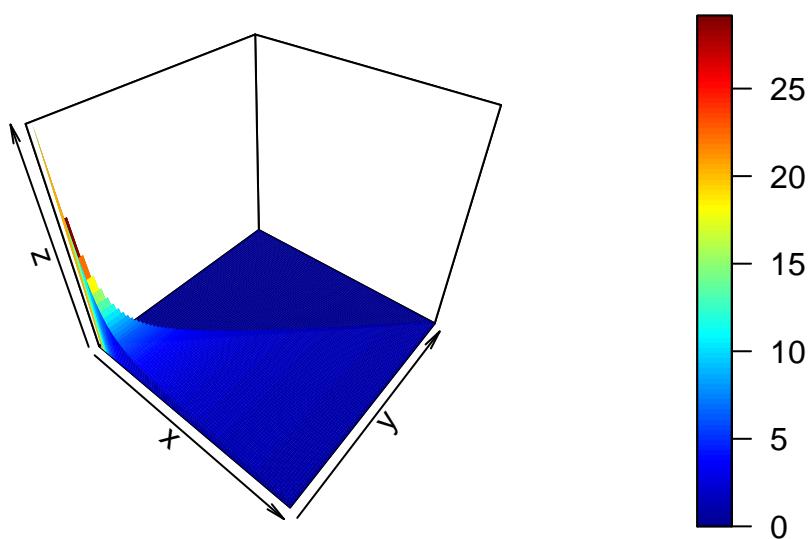


Figure 3: Approximate joint pdf of  $X$  and  $Y$

**4. Find the conditional probability density function of  $X$  given  $Y$ ,  $f_{X|Y}$ .**

In order to find  $f_{X|Y}$  we need the marginal pdf of  $Y$ .

$$f_Y(y) = \int_{\mathbb{X}} f_{X,Y}(x,y) \cdot dx = \int_{y=0}^x \frac{1}{x} \cdot dx = \int_{x=y}^1 \frac{dx}{x} = [\log(x)]_{x=y}^1 dx = -\log(y) = \log\left(\frac{1}{y}\right)$$

This result confirms what the shape of  $f_Y(y)$  in Figure 1 suggested.

```
# f(y)
pdf_y <- function(y)
  sapply(y, function(y) integrate(function(x)
    pdf_y_given_x(x,y)*pdf_x(x), 0, 1)$value)
integrate(pdf_y, -Inf, Inf) # integral

## 1 with absolute error < 9.3e-05

plot(sort(y), pdf_y(sort(y)), type = 'l', main = 'pdf(y)', xlab = 'y')
# Confirm that f(y) = log(1/y)
lines(sort(y), log(1/sort(y)), type = 'l', main = 'pdf(y)')
```

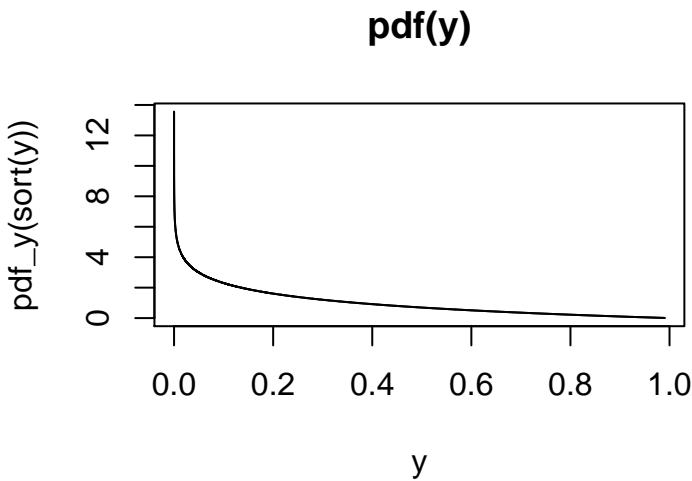


Figure 4: Approximate pdf of  $Y$  conditional on  $X$  for two values of  $X$

$$f_{X|Y}(x|y) = \frac{f_{X,Y}(x,y)}{f_Y(y)} = \frac{-1}{x \cdot \log(y)}$$

5. Find the expectation of  $X$ , given that  $Y$  is  $1/2$ ,  $E(X|Y = 1/2)$ .

$$\begin{aligned} E(X|Y = 1/2) &= \int_{\mathbb{X}} x \cdot f_{X|Y}(x|y = 1/2) \cdot dx = \int_{x=1/2}^1 x \cdot \left( \frac{-1}{x \cdot \log(1/2)} \right) \cdot dx \\ &= \frac{1}{\log(2)} \int_{x=1/2}^1 dx = \frac{1}{\log(2)} [x]_{x=1/2}^1 = \frac{1}{2 \cdot \log(2)} = 0.721 \end{aligned}$$

```
# Confirm E(X/Y=0.5) (use values of Y around 0.5 in the previous simulation)
mean(x[y > 0.5 - 1e-2 & y < 0.5 + 1e-2])
```

```
## [1] 0.72847
```

```
1/(2*log(2))
```

```
## [1] 0.7213475
```

## Question 2: Investing

Suppose that you are planning an investment in three different companies. The payoff per unit you invest in each company is represented by a random variable. A represents the payoff per unit invested in the first company, B in the second, and C in the third. A, B, and C are independent of each other. Furthermore,  $\text{Var}(A) = 2\text{Var}(B) = 3\text{Var}(C)$ .

You plan to invest a total of one unit in all three companies. You will invest amount a in the first company, b in the second, and c in the third, where  $a, b, c \in [0, 1]$  and  $a + b + c = 1$ . Find, the values of a, b, and c that minimize the variance of your total payoff.

Let's call P the total payoff:

$$\text{Var}(P) = \text{Var}(aA + bB + cC) = a^2\text{Var}(A) + b^2\text{Var}(B) + c^2\text{Var}(C)$$

because A, B, and C are independent of each other. And since  $\text{Var}(A) = 2\text{Var}(B) = 3\text{Var}(C)$ , we can derive that:

$$\text{Var}(P) = \text{Var}(A) \left( a^2 + \frac{b^2}{2} + \frac{c^2}{3} \right)$$

We want to:

$$\begin{array}{ll} \text{minimize} & P(a, b, c) \\ \text{subject to} & g(a, b, c) = 0 \end{array}$$

where  $g(a, b, c) = a + b + c - 1$ , so  $g(a, b, c) = 0$  is our constraint.

Using the Lagrange multiplier method, we can define:

$$\mathcal{L}(a, b, c, \lambda) = P(a, b, c) - \lambda \cdot g(a, b, c)$$

So to find our local minima we need to solve:

$$\nabla_{a,b,c,\lambda} \mathcal{L} = 0$$

$$\left( \frac{\partial \mathcal{L}}{\partial a}, \frac{\partial \mathcal{L}}{\partial b}, \frac{\partial \mathcal{L}}{\partial c}, \frac{\partial \mathcal{L}}{\partial \lambda} \right) = \left( 2a - \lambda, b - \lambda, \frac{2}{3}c - \lambda, -(a + b + c - 1) \right) = \mathbf{0}$$

$$\Rightarrow \begin{cases} 2a - \lambda = 0 \\ b - \lambda = 0 \\ 2c/3 - \lambda = 0 \\ a + b + c - 1 = 0 \end{cases} \Rightarrow \begin{cases} a = \lambda/2 \\ b = \lambda \\ c = 3\lambda/2 \\ \frac{\lambda}{2} + \lambda + \frac{3\lambda}{2} = 3\lambda = 1 \end{cases} \Rightarrow \begin{cases} \mathbf{a = \frac{1}{6} = 0.1667} \\ \mathbf{b = \frac{1}{3} = 0.3333} \\ \mathbf{c = \frac{1}{2} = 0.5} \end{cases}$$

Let's prove the result in R:

```
payoff <- function(x) {
  a <- x[1]
  b <- x[2]
  c <- x[3]
  a^2 + b^2/2 + c^2/3
}
gradient_payoff <- function(x) {
  a <- x[1]
  b <- x[2]
  c <- x[3]
  c(2*a, b, 2*c/3)
}
sol <- constrOptim(theta = c(.3, .3, .4), f = payoff, grad = gradient_payoff,
                    ui = rbind(c(1, 0, 0), c(0, 1, 0), c(0, 0, 1),
                               c(-1, 0, 0), c(0, -1, 0), c(0, 0, -1),
                               c(1, 1, 1), c(-1, -1, -1)),
                    ci = c(0, 0, -1, -1, -1, 1-1e-6, -1-1e-6))
sol$par
```

```
## [1] 0.1666989 0.3333673 0.4999327
```

## Question 3: Turtles

Next, suppose that the lifespan of a species of turtle follows a uniform distribution over  $[0, \theta]$ . Here, parameter  $\theta$  represents the unknown maximum lifespan. You have a random sample of  $n$  individuals, and measure the lifespan of each individual  $i$  to be  $y_i$ .

1. Write down the likelihood function,  $l(\theta)$  in terms of  $y_1, y_2, \dots, y_n$ .

$$l(\theta; y_1, \dots, y_n) = f(y_1, \dots, y_n | \theta) = \prod_{i=1}^n f(y_i | \theta) = \begin{cases} \prod_{i=1}^n \frac{1}{\theta} = \theta^{-n} & 0 \leq y_i \leq \theta \\ 0 & \text{otherwise} \end{cases}$$

2. Based on the previous result, what is the maximum-likelihood estimator for  $\theta$ ?

The MLE of  $\theta$  must be a value of  $\theta$  for which  $\theta \geq y_i$  for  $i = 1, \dots, n$  and which maximizes  $1/\theta^n$  among all such values. I.e., the maximum value of  $y_i$  within the sample.

$$\hat{\theta}_{ml} = \arg \max_{\theta \in \Theta} \hat{l}(\theta; y_1, \dots, y_n) = \max\{y_1, \dots, y_n\}$$

3. Let  $\hat{\theta}_{ml}$  be the maximum likelihood estimator above. For the simple case that  $n = 1$ , what is the expectation of  $\hat{\theta}_{ml}$ , given  $\theta$ ?

$$E(\hat{\theta}_{ml} | \theta) = E(y_1) = E(y) = \int_{y=0}^{\theta} \frac{y}{\theta} \cdot dy = \left[ \frac{y^2}{2\theta} \right]_{y=0}^{\theta} = \frac{\theta}{2}$$

4. Is the maximum likelihood estimator biased?

Yes, it is:

$$E(\hat{\theta}_{ml}) - \theta = \frac{\theta}{2} \neq 0$$

5. For the more general case that  $n \geq 1$ , what is the expectation of  $\hat{\theta}_{ml}$ ?

Without loss of generality, let's suppose the individual  $n$  is the one with the maximum lifespan among the sample, i.e.,  $y_n \geq y_i$  for  $i = 1, \dots, n-1$ . Call  $z$  that maximum value of  $y_i$ .

$$E[\max\{y_1, \dots, y_n\}] = E(y_n) = E(z) = \int_{z=0}^{\theta} z \cdot f(z) dz$$

But what is the distribution of  $z$ ?

$$\begin{aligned} F(z) &= \Pr(y_n \leq z) = \Pr(y_1 \leq z \cap \dots \cap y_n \leq z) \stackrel{\text{i.i.d.}}{=} \prod_{i=1}^n \Pr(y_i \leq z) = \left( \frac{z}{\theta} \right)^n \Rightarrow f(z) = \frac{n z^{n-1}}{\theta^n} \\ E[\max\{y_1, \dots, y_n\}] &= \frac{n}{\theta^n} \int_{z=0}^{\theta} z^n dz = \frac{n}{\theta^n} \left[ \frac{z^{n+1}}{n+1} \right]_{z=0}^{\theta} = \frac{n}{n+1} \cdot \theta \end{aligned}$$

(Which confirms the previous result, for  $n = 1$ .)

### 6. Is the maximum likelihood estimator consistent?

It is:

$$\Pr \left( |\hat{\theta}_{ml} - \theta| > \varepsilon \right) = \Pr \left( \frac{\theta}{n+1} > \varepsilon \right) \rightarrow 0 \text{ as } n \rightarrow \infty$$

```
simulations <- 1e3 # number of simulations
theta <- 100 # an arbitrary value of theta
y <- runif(n = simulations, min = 0, max = theta) # Y ~ U(0,theta)
# any(y == theta); all(y < theta) # FALSE and TRUE, respectively
# No matter how large is the sample, Yi is always lower than 1
set.seed(1)
num_simulations <- sort(c(1, sample(c(2:simulations), 49)))
theta_mle <- unlist(lapply(num_simulations, function(n)
  mean(max(runif(n = n, min = 0, max = theta)))))
plot(num_simulations, theta_mle, ylim = c(floor(min(theta_mle)), theta),
     xlab = "Number of simulations", ylab = "MLE of theta", pch = '*')
lines(num_simulations, theta_mle, lwd = 0.5, col = 'blue')
```

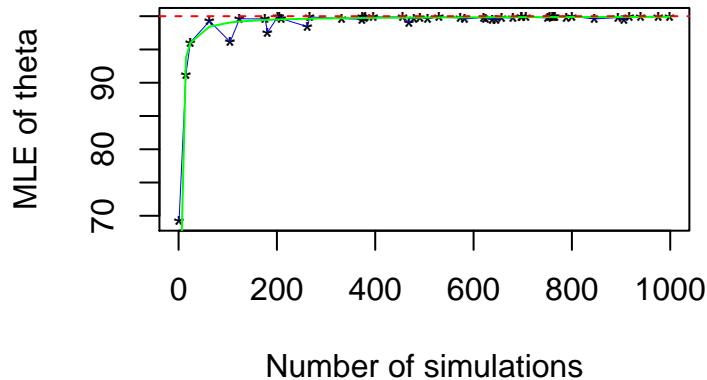


Figure 5: Trend of the MLE of  $\theta$  depending on the sample size

## Question 4: CLM 1

### Background

The file `WageData2.csv` contains a dataset that has been used to quantify the impact of education on wage. One of the reasons we are proving another wage-equation exercise is that this area by far has the most (and most well-known) applications of instrumental variable techniques, the endogeneity problem is obvious in this context, and the datasets are easy to obtain.

### The Data

You are given a sample of 1000 individuals with their wage, education level, age, working experience, race (as an indicator), father's and mother's education level, whether the person lived in a rural area, whether the person lived in a city, IQ score, and two potential instruments, called `z1` and `z2`.

The dependent variable of interest is wage (or its transformation), and we are interested in measuring “return” to education, where return is measured in the increase (hopefully) in wage with an additional year of education.

### Question 4.1

Conduct an univariate analysis (using tables, graphs, and descriptive statistics found in the last 7 lectures) of all of the variables in the dataset.

Also, create two variables: (1) natural log of wage (name it `logWage`) (2) square of experience (name it `experienceSquare`)

Table 1: Summary Statistics of Wage Data

Statistic	N	Mean	St. Dev.	Min	Pctl(25)	Median	Pctl(75)	Max
X	1,000	1,466.678	866.535	5	715.5	1,431.5	2,212	3,009
wage	1,000	578.783	266.569	127	400	543	702.5	2,404
education	1,000	13.219	2.729	2	12	12	16	18
experience	1,000	8.788	4.221	0	6	8	11	23
age	1,000	28.007	3.118	24	25	27	30	34
raceColor	1,000	0.238	0.426	0	0	0	0	1
dad_education	761	10.181	3.748	0	8	11	12	18
mom_education	872	10.451	3.126	0	8	12	12	18
rural	1,000	0.391	0.488	0	0	0	1	1
city	1,000	0.712	0.453	0	0	1	1	1
z1	1,000	0.440	0.497	0	0	0	1	1
z2	1,000	0.686	0.464	0	0	1	1	1
IQscore	684	102.273	15.843	50	93	103	113	144
logWage	1,000	6.263	0.447	4.844	5.991	6.297	6.555	7.785
experienceSquare	1,000	95.030	86.786	0	36	64	121	529

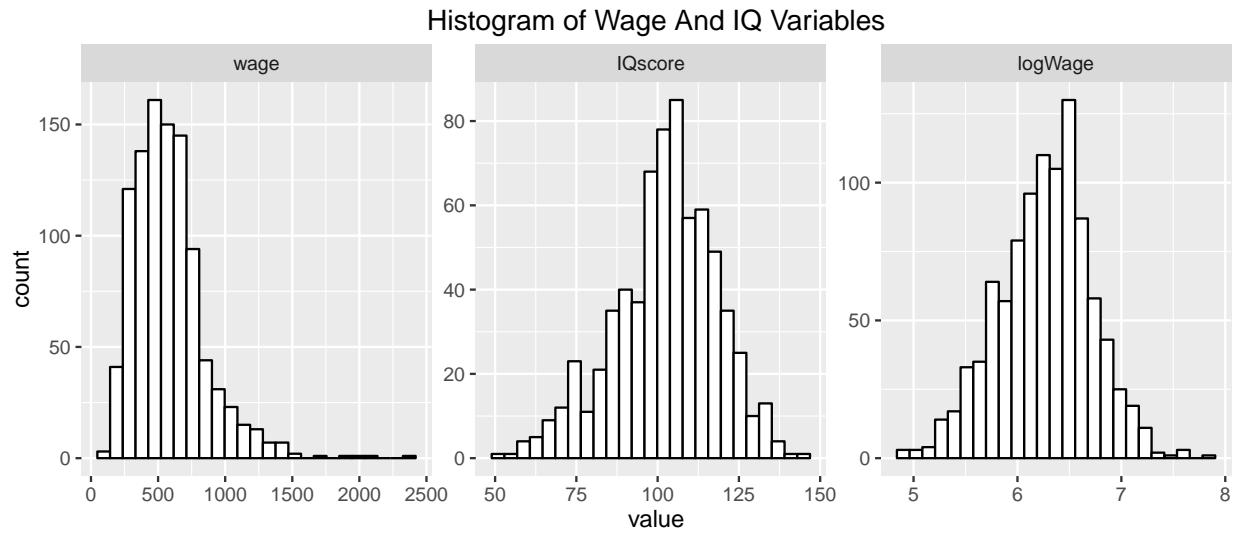


Figure 6: Histogram of Wage , IQ Score, and log(Wage)

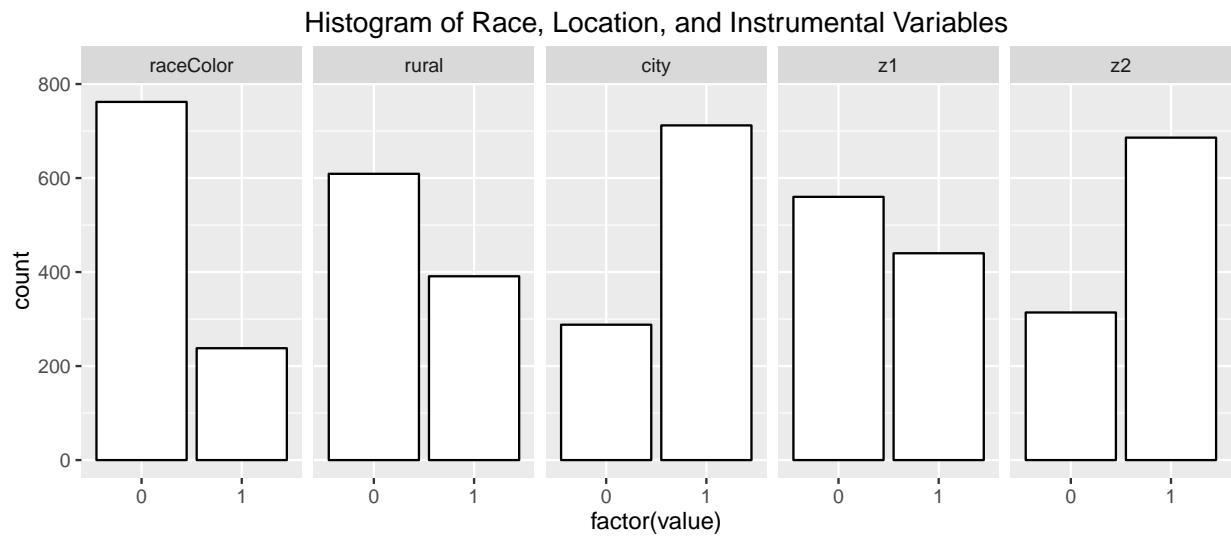


Figure 7: Histogram of Race , Rural, City, Z1, and Z2

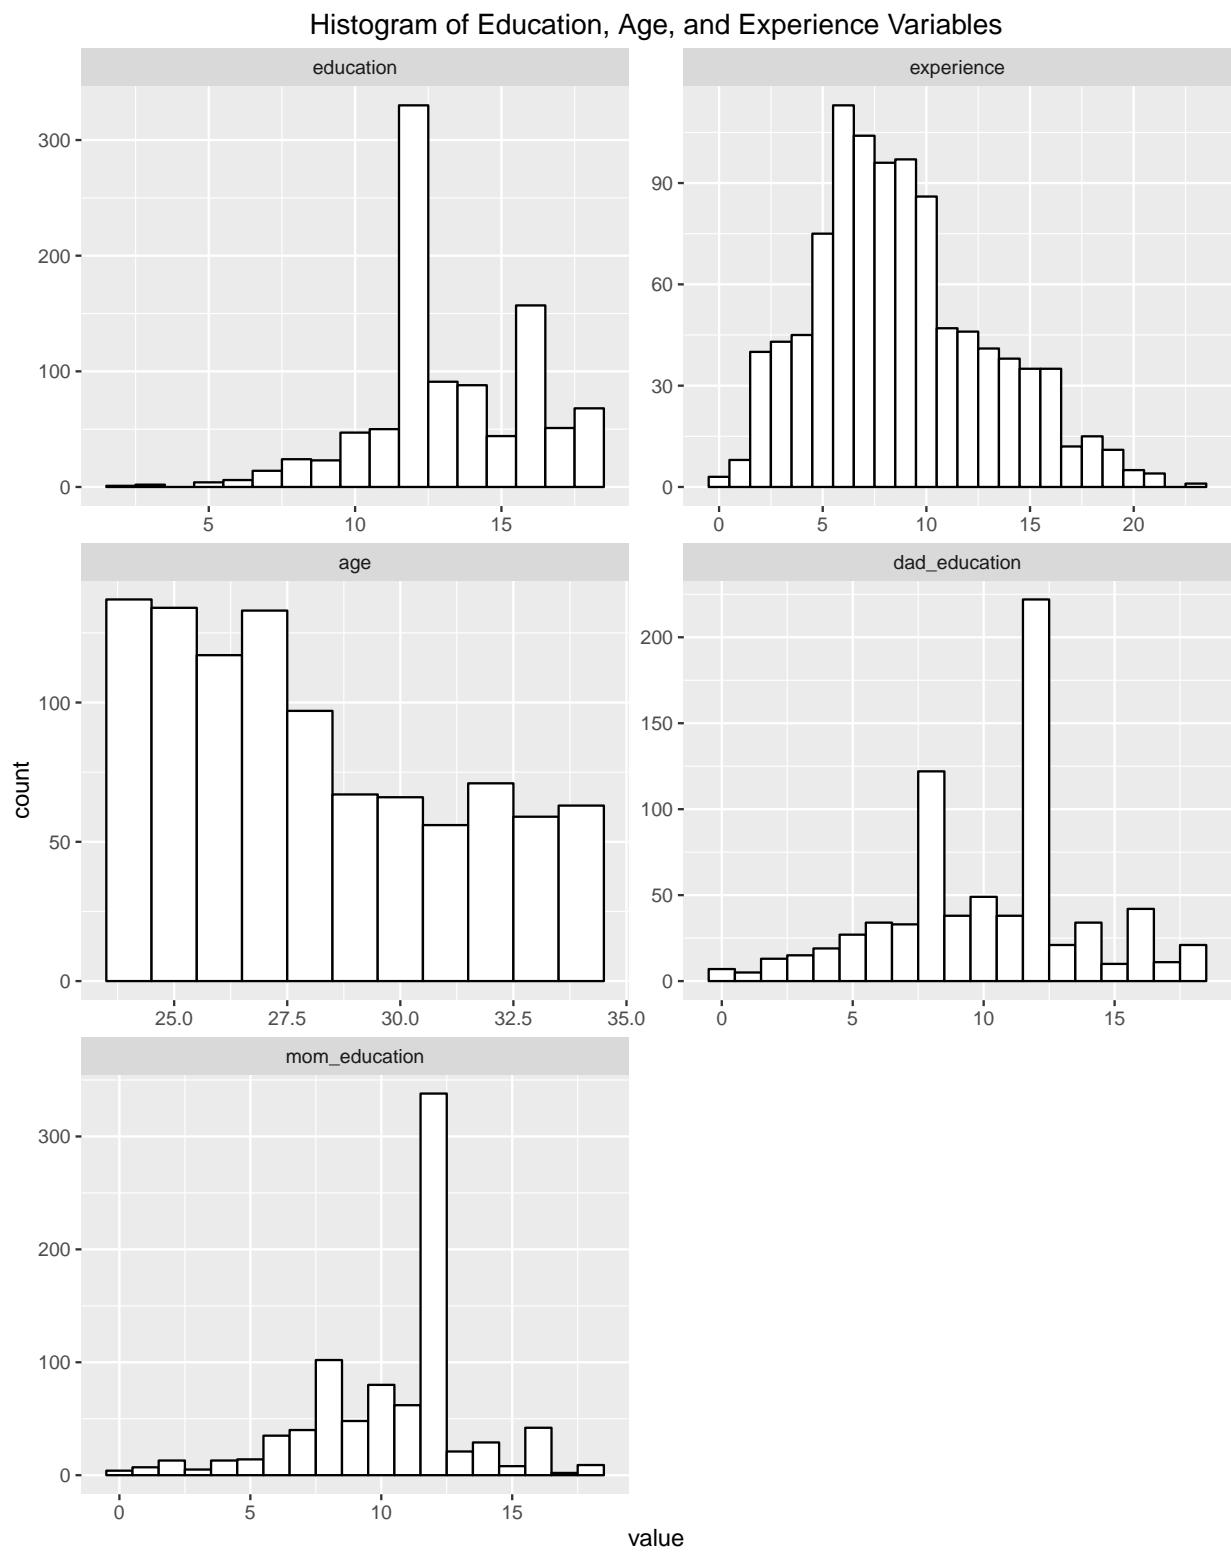


Figure 8: Histogram of Education, Experience, Age, Father’s Education, and Mother’s Education

**Question 4.2**

Conduct a bivariate analysis (using tables, graphs, descriptive statistics found in the last 7 lectures) of wage and logWage and all the other variables in the datasets.

Table 2: Correlations for Wage and log(Wage)

	Wage	log(Wage)
wage	1	0.946
education	0.310	0.332
experience	-0.006	-0.029
age	0.264	0.251
raceColor	-0.301	-0.341
dad_education	0.190	0.189
mom_education	0.198	0.210
rural	-0.222	-0.250
city	0.220	0.236
z1	0.101	0.087
z2	0.171	0.177
IQscore	0.186	0.201
logWage	0.946	1
experienceSquare	-0.043	-0.065

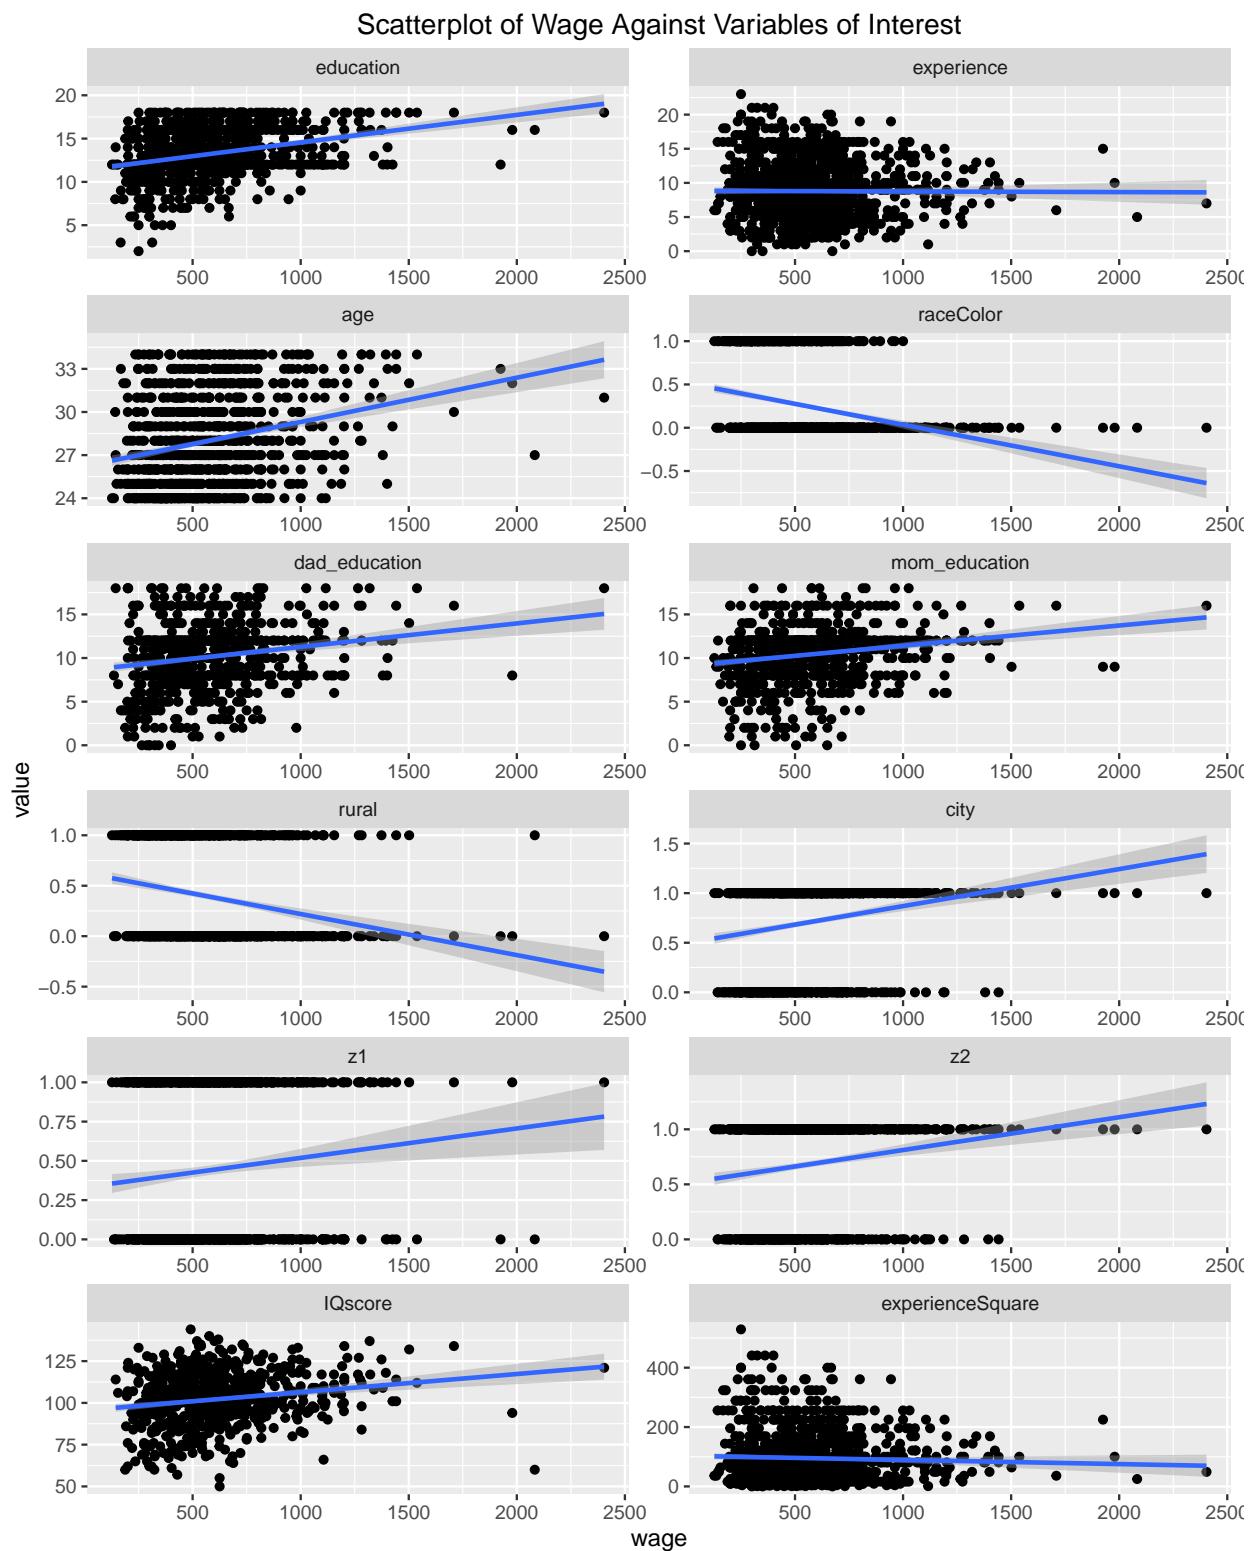


Figure 9: Scatterplot of Wage Against Variables of Interest

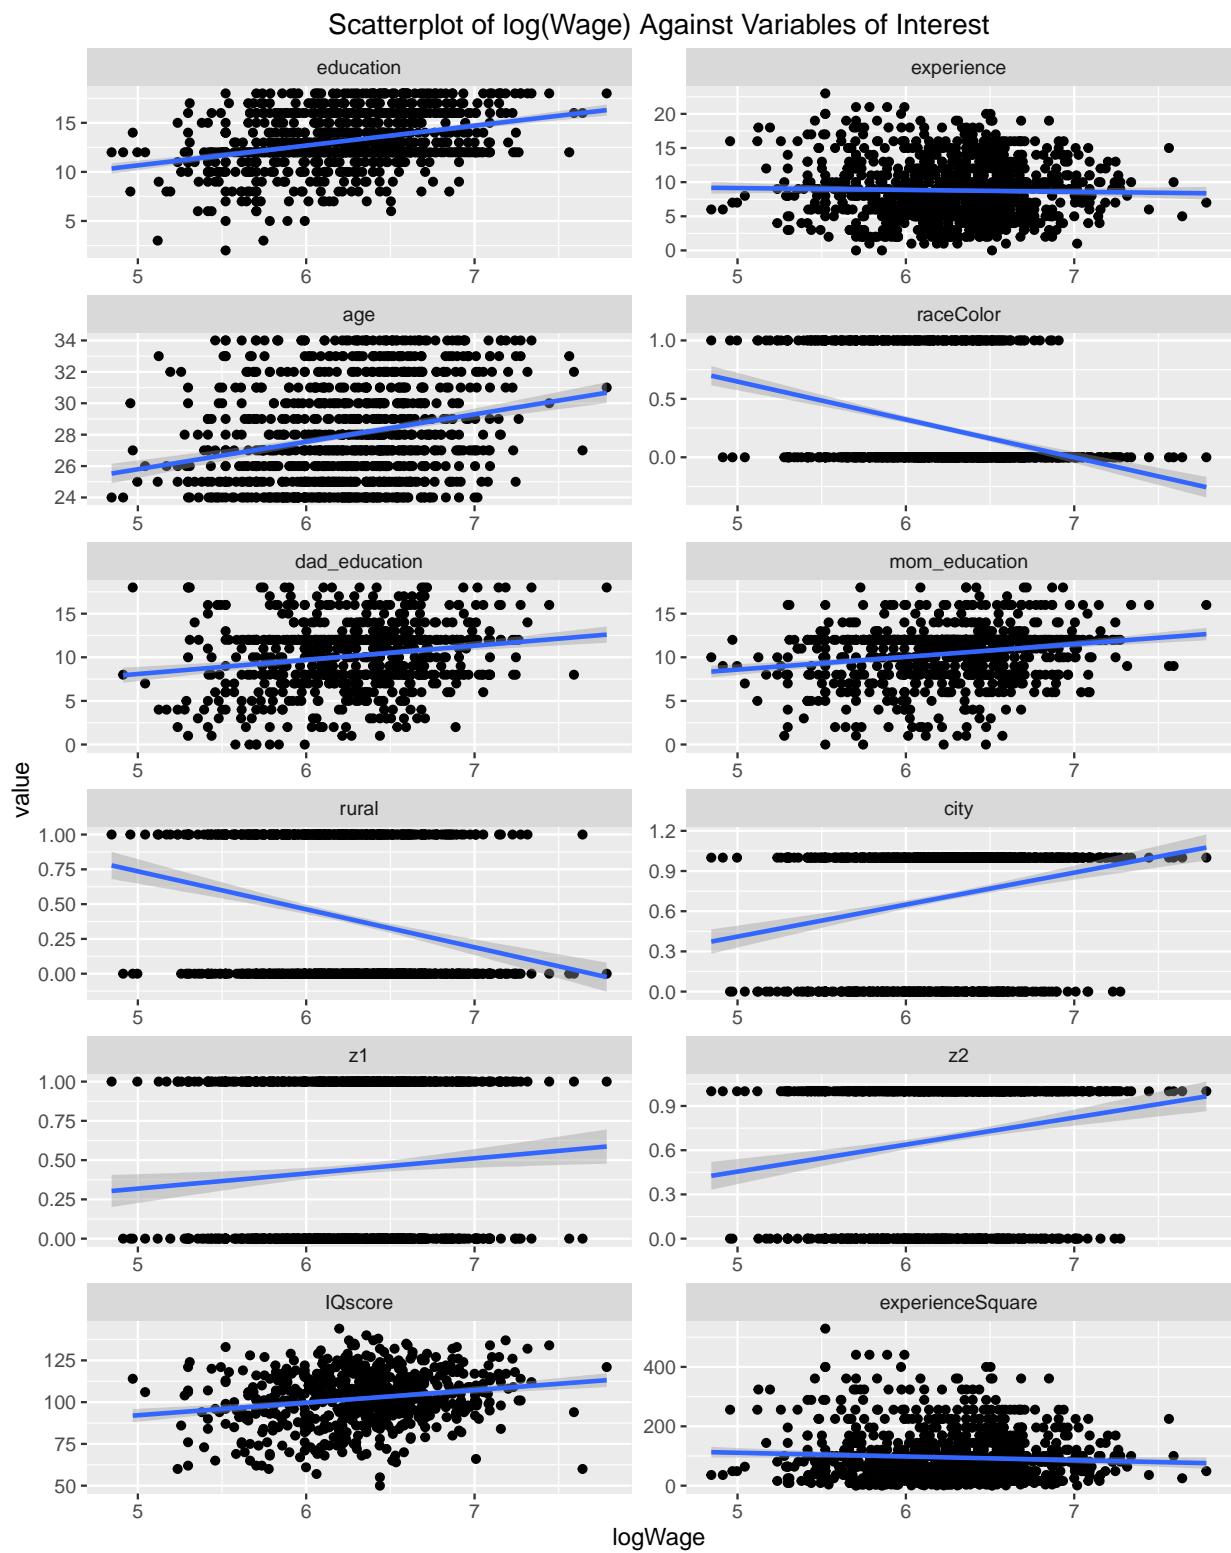


Figure 10: Scatterplot of log(Wage) Against Variables of Interest

### Question 4.3

Regress  $\log(wage)$  on education, experience, age, and raceColor.

- Report all the estimated coefficients, their standard errors, t-statistics, F-statistic of the regression,  $R^2$ ,  $R^2_{adj}$ , and degrees of freedom.

Table 3: Regression Summary

<i>Dependent variable:</i>	
	logWage
education	0.080*** (0.006) t = 12.486
experience	0.035*** (0.004) t = 8.869
age	
raceColor	-0.261*** (0.030) t = -8.564
Constant	4.962*** (0.113) t = 43.774
Observations	1,000
$R^2$	0.236
Adjusted $R^2$	0.234
Residual Std. Error	0.392 (df = 996)
F Statistic	102.582*** (df = 3; 996)

Note: \*p<0.1; \*\*p<0.05; \*\*\*p<0.01

- Explain why the degrees of freedom takes on the specific value you observe in the regression output.

The overall degrees of freedom (F(3, 996) is one smaller than we might expect otherwise. This is because the parameter for age is a linearly dependent combination of the other parameters and its coefficient cannot be estimated. R automatically removes this variable from the model, and thus our regression output reflects the degrees of freedom for a model with 3 parameters rather than 4. Thus estimating this model is equivalent to estimating  $\log(wage) = education + experience + raceColor$ .

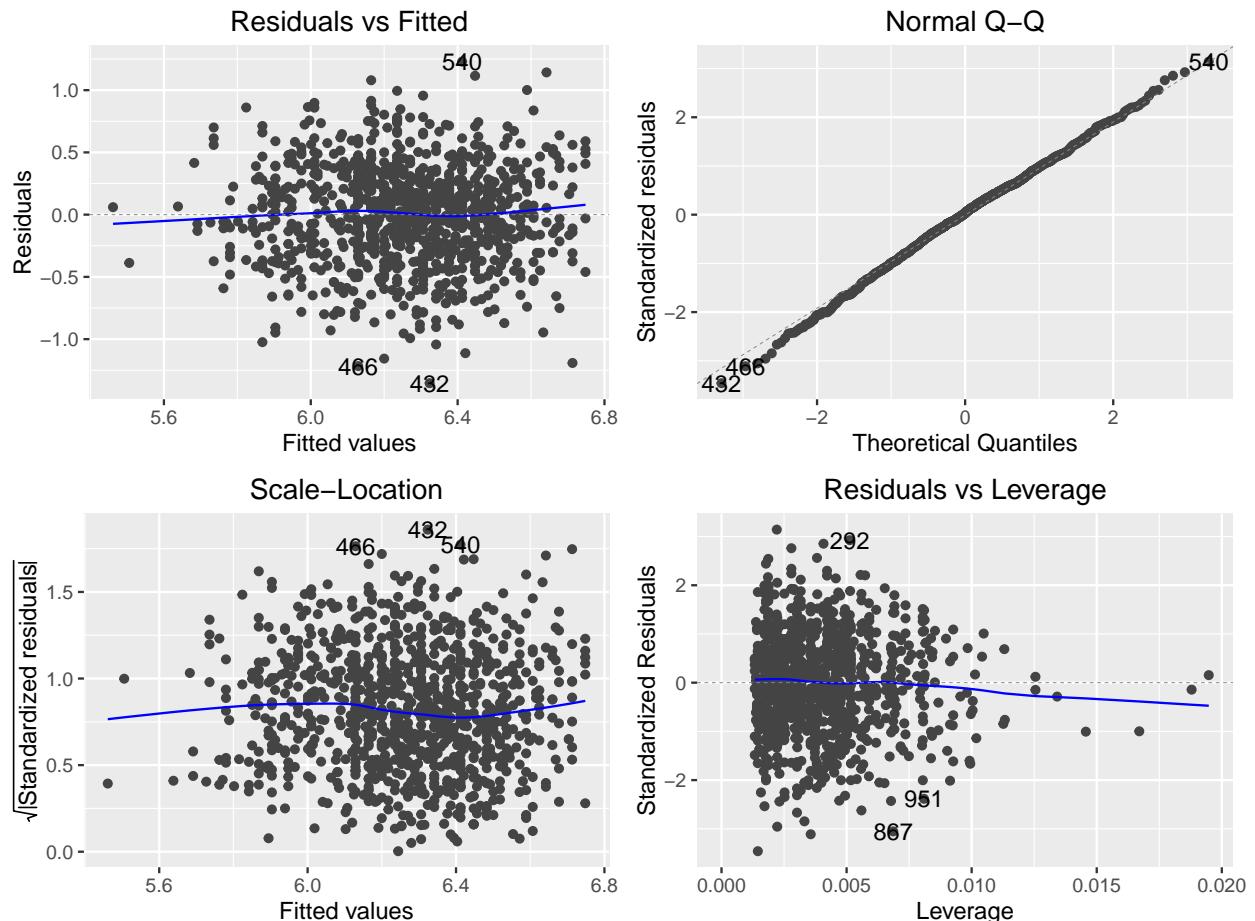


Figure 11: Regression Diagnostic Plots for Question 4.3

Table 4: Regression summary

<i>Dependent variable:</i>	
	log(Wages)
Education	0.080*** (0.006)
Experience	0.035*** (0.004)
Race (White or Non-white)	-0.261*** (0.030)
Constant	4.962*** (0.115)
F Statistic	100.278***
df	3; 996
Observations	1,000
R <sup>2</sup>	0.236
Adjusted R <sup>2</sup>	0.234
Residual Std. Error	0.392

·p<0.1; \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

3. Describe any unexpected results from your regression and how you would resolve them (if the intent is to estimate return to education, condition on race and experience).

The inability to estimate the age coefficient is an unexpected result. Looking at the correlation matrix from above, we don't see that any variable is perfectly correlated with  $\log(wage)$ . However, if we take the correlation of  $age$  with the sum of  $education$  and  $experience$ , we see that  $age$  is perfectly correlated with the two variables. Since the model already incorporates all of the information contained in the age variable, we can simply remove it from the model and not lose anything.

#### 4. Interpret the coefficient estimate associated with education.

Holding other covariates constant, a one year increase in  $education$  was associated with a statistically significant increase in  $\log(wage)$ , ( $\beta = 0.080$ , std. error = 0.006,  $t = 12.394$ ,  $p < .001$ ). This effect also represents a practically significant wage return on education.

#### 5. Interpret the coefficient estimate associated with experience.

Holding other covariates constant, a one year increase in  $experience$  was associated with a statistically significant increase in  $\log(wage)$ , ( $\beta = 0.035$ , std. error = 0.004,  $t = 8.81$ ,  $p < .001$ ). This effect also represents a practically significant wage return on experience.

### Question 4.4

Regress  $\log(wage)$  on education, experience, experienceSquare, and race-Color.

1. Plot a graph of the estimated effect of experience on wage.

2. What is the estimated effect of experience on wage when experience is 10 years?

The estimated effect of experience at 10 years on wages is  $10 \cdot 0.092 + 10^2 \cdot -0.003 = 0.637$

Table 5: Regression summary

<i>Dependent variable:</i>	
	log(Wages)
Education	0.079*** (0.006)
Experience	0.092*** (0.011)
Expereince <sup>2</sup>	-0.003*** (0.001)
Race (White or Non-white)	-0.263*** (0.030)
Constant	4.736*** (0.120)
F Statistic	84.960***
df	4; 995
Observations	1,000
R <sup>2</sup>	0.257
Adjusted R <sup>2</sup>	0.254
Residual Std. Error	0.387

·p<0.1; \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

## Question 4.5

Regress logWage on education, experience, experienceSquare, raceColor, dad\_education, mom\_education, rural, city.

1. What are the number of observations used in this regression? Are missing values a problem? Analyze the missing values, if any, and see if there is any discernible pattern with wage, education, experience, and raceColor.

There are 723 observations in the regression. A comparison of the original data and the observations in the regression reveals that the missing observations tend to have lower wages and education and are less likely to live in the city. The missing observations have more experience, and are more likely to be non-white and live in rural areas.

2. Do you just want to “throw away” these observations?

The clear pattern of missing values representing more rural, non-white observations with lower education indicates that these values should not simply be discarded from the model. In order to accurately assess the effect of covariates these values are important for the regression.

3. How about blindly replace all of the missing values with the average of the observed values of the corresponding variable? Rerun the original regression using all of the observations?

```
data7 <- data
data7$dad_education <- na.fill(data7$dad_education, mean(data7$dad_education,
                                         na.rm=T))
data7$mom_education <- na.fill(data7$mom_education, mean(data7$mom_education,
```

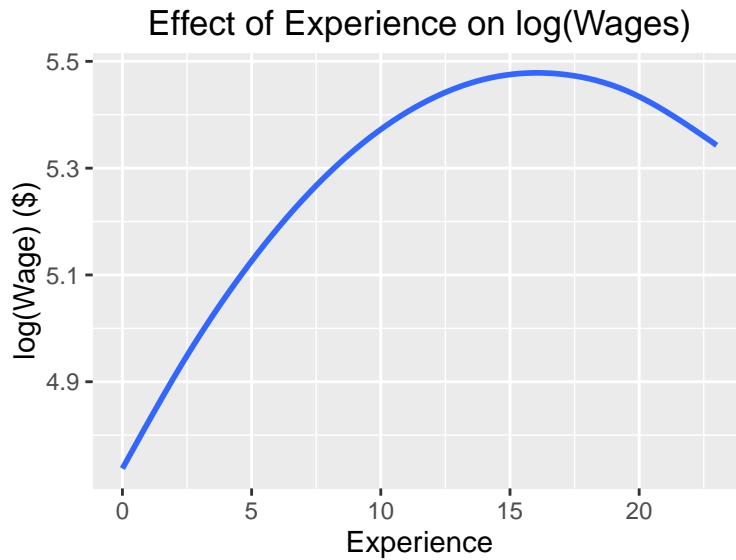


Figure 12: Estimated Effect of Experience on Wages

Table 6: Regression summary

<i>Dependent variable:</i>	
	log(Wages)
Education	0.068*** (0.008)
Experience	0.097*** (0.013)
Expereince <sup>2</sup>	-0.003*** (0.001)
Race (White or Non-white)	-0.213*** (0.041)
Father's Education	-0.001 (0.006)
Mother's Education	0.011 (0.007)
Rural (Yes or No)	-0.092** (0.032)
City (Yes or No)	0.178*** (0.032)
Constant	4.642*** (0.150)
F Statistic	32.592***
df	8; 714
Observations	723
R <sup>2</sup>	0.275
Adjusted R <sup>2</sup>	0.267
Residual Std. Error	0.379

·p<0.1; \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

Table 7: All Observations

Statistic	N	Mean	St. Dev.	Min	Pctl(25)	Median	Pctl(75)	Max
X	1,000	1,466.678	866.535	5	715.5	1,431.5	2,212	3,009
wage	1,000	578.783	266.569	127	400	543	702.5	2,404
education	1,000	13.219	2.729	2	12	12	16	18
experience	1,000	8.788	4.221	0	6	8	11	23
age	1,000	28.007	3.118	24	25	27	30	34
raceColor	1,000	0.238	0.426	0	0	0	0	1
dad_education	761	10.181	3.748	0	8	11	12	18
mom_education	872	10.451	3.126	0	8	12	12	18
rural	1,000	0.391	0.488	0	0	0	1	1
city	1,000	0.712	0.453	0	0	1	1	1
z1	1,000	0.440	0.497	0	0	0	1	1
z2	1,000	0.686	0.464	0	0	1	1	1
IQscore	684	102.273	15.843	50	93	103	113	144
logWage	1,000	6.263	0.447	4.844	5.991	6.297	6.555	7.785
experienceSquare	1,000	95.030	86.786	0	36	64	121	529

Table 8: Missing Father's Education

Statistic	N	Mean	St. Dev.	Min	Pctl(25)	Median	Pctl(75)	Max
X	239	1,675.259	864.761	15	921.5	1,712	2,515.5	3,009
wage	239	529.531	267.216	127	351.5	465	640.5	2,083
education	239	12.121	2.696	3	11	12	13	18
experience	239	10.527	4.223	1	7	10	14	23
age	239	28.649	3.090	24	26	28	31	34
raceColor	239	0.464	0.500	0	0	0	1	1
mom_education	149	8.987	3.397	0	7	9	12	18
rural	239	0.540	0.499	0	0	1	1	1
city	239	0.661	0.474	0	0	1	1	1
z1	239	0.397	0.490	0	0	0	1	1
z2	239	0.669	0.471	0	0	1	1	1
IQscore	132	96.038	17.421	50	85	98.5	107	135
logWage	239	6.163	0.465	4.844	5.862	6.142	6.462	7.642
experienceSquare	239	128.577	98.151	1	49	100	196	529

Table 9: Missing Mother's Education

Statistic	N	Mean	St. Dev.	Min	Pctl(25)	Median	Pctl(75)	Max
X	128	1,765.523	870.253	53	996	1,796	2,572	3,009
wage	128	525.938	252.627	142	371.5	492.5	625	2,083
education	128	11.891	2.713	2	10	12	13	18
experience	128	10.898	4.742	0	7	10	15	21
age	128	28.789	3.513	24	26	28	32	34
raceColor	128	0.461	0.500	0	0	0	1	1
dad_education	38	9.184	3.840	2	6	9.5	12	16
rural	128	0.523	0.501	0	0	1	1	1
city	128	0.656	0.477	0	0	1	1	1
z1	128	0.414	0.494	0	0	0	1	1
z2	128	0.680	0.468	0	0	1	1	1
IQscore	71	93.352	15.900	60	80.5	96	104.5	124
logWage	128	6.175	0.418	4.956	5.918	6.199	6.438	7.642
experienceSquare	128	141.086	108.986	0	49	100	225	441

Table 10: Actual Observations

Statistic	N	Mean	St. Dev.	Min	Pctl(25)	Median	Pctl(75)	Max
X	723	1,399.028	853.489	5	680.5	1,314	2,125	2,998
wage	723	597.080	268.094	136	409	570	721	2,404
education	723	13.651	2.616	3	12	13	16	18
experience	723	8.145	4.005	0	5	8	10	21
age	723	27.797	3.066	24	25	27	30	34
raceColor	723	0.158	0.365	0	0	0	0	1
dad_education	723	10.234	3.739	0	8	11	12	18
mom_education	723	10.752	2.982	0	9	12	12	18
rural	723	0.346	0.476	0	0	0	1	1
city	723	0.730	0.444	0	0	1	1	1
z1	723	0.448	0.498	0	0	0	1	1
z2	723	0.685	0.465	0	0	1	1	1
IQscore	531	104.081	15.062	60	95	105	114	144
logWage	723	6.297	0.442	4.913	6.014	6.346	6.581	7.785
experienceSquare	723	82.367	78.067	0	25	64	100	441

```
na.rm=T))
model6 <- lm(logWage ~ education + experience + experienceSquare + raceColor +
               dad_education + mom_education + rural + city, data = data7)
```

Table 11: Regression summary

<i>Dependent variable:</i>	
	log(Wages)
Education	0.071*** (0.007)
Experience	0.090*** (0.011)
Expereince <sup>2</sup>	-0.003*** (0.001)
Race (White or Non-white)	-0.231*** (0.031)
Father's Education	-0.00004 (0.005)
Mother's Education	0.003 (0.005)
Rural (Yes or No)	-0.095*** (0.027)
City (Yes or No)	0.167*** (0.026)
Constant	4.729*** (0.128)
F Statistic	54.109***
df	8; 991
Observations	1,000
R <sup>2</sup>	0.298
Adjusted R <sup>2</sup>	0.292
Residual Std. Error	0.376

.p<0.1; \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

4. How about regress the variable(s) with missing values on education, experience, and raceColor, and use this regression(s) to predict (i.e., “impute”) the missing values and then rerun the original regression using all of the observations?

```
data8 <- data
model_dad<-lm(dad_education~education+experience+raceColor, data=data8)
model_mom<-lm(mom_education~education+experience+raceColor, data=data8)
coef_dad<-coef(model_dad)
coef_mom<-coef(model_mom)
# Impute values for missing observations
for (i in 1:nrow(data8)) {
  if (is.na(data8$dad_education[i])==TRUE) {
    data8$dad_education[i]= coef_dad[1]+coef_dad[2]*data8$education[i]+
      coef_dad[3]*data8$experience[i]+coef_dad[4]*data8$raceColor[i]
  }
  if (is.na(data8$mom_education[i])==TRUE) {
```

```

    data8$mom_education[i] = coef_mom[1]+coef_mom[2]*data8$education[i] +
      coef_mom[3]*data8$experience[i]+coef_mom[4]*data8$raceColor[i]
  }
}

```

Table 12: Regression summary

<i>Dependent variable:</i>	
	log(Wages)
Education	0.070*** (0.007)
Experience	0.089*** (0.011)
Expereince <sup>2</sup>	-0.003*** (0.001)
Race (White or Non-white)	-0.227*** (0.031)
Father's Education	0.002 (0.005)
Mother's Education	0.002 (0.006)
Rural (Yes or No)	-0.095*** (0.027)
City (Yes or No)	0.167*** (0.026)
Constant	4.727*** (0.125)
F Statistic	54.147***
df	8; 991
Observations	1,000
R <sup>2</sup>	0.298
Adjusted R <sup>2</sup>	0.293
Residual Std. Error	0.376

·p<0.1; \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

5. Compare the results of all of these regressions. Which one, if at all, would you prefer?

Table 13:

	Dependent variable:		
	(1)	(2)	(3)
Education	0.068*** (0.008)	0.071*** (0.007)	0.070*** (0.007)
Experience	0.097*** (0.013)	0.090*** (0.011)	0.089*** (0.011)
Expereince <sup>2</sup>	-0.003*** (0.001)	-0.003*** (0.001)	-0.003*** (0.001)
Race (White or Non-white)	-0.213*** (0.041)	-0.231*** (0.031)	-0.227*** (0.031)
Father's Education	-0.001 (0.006)	-0.00004 (0.005)	0.002 (0.005)
Mother's Education	0.011 (0.007)	0.003 (0.005)	0.002 (0.006)
Rural (Yes or No)	-0.092** (0.032)	-0.095*** (0.027)	-0.095*** (0.027)
City (Yes or No)	0.178*** (0.032)	0.167*** (0.026)	0.167*** (0.026)
Constant	4.642*** (0.150)	4.729*** (0.128)	4.727*** (0.125)
F Statistic	32.592***	54.109***	54.147***
df	8; 714	8; 991	8; 991
Observations	723	1,000	1,000
R <sup>2</sup>	0.275	0.298	0.298
Adjusted R <sup>2</sup>	0.267	0.292	0.293
Residual Std. Error	0.379	0.376	0.376

·p<0.1; \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

Comparing the regression results, it's worth noting that the magnitude and significance of the coefficients are broadly similar. However, given that the values that we imputed are for covariates that are neither statistically nor practically significant, I would prefer a model that used all of the observations and did not include the covariates for parent's education.

## Question 4.6

1. Consider using  $z_1$  as the instrumental variable (IV) for education. What assumptions are needed on  $z_1$  and the error term (call it,  $u$ )?

Formally, we need to assume that the instrumental variable is uncorrelated with the error term. Ie.

$$\text{Cov}(z_{1j}, u) = 0 \quad \forall j = 1, 2, \dots, k$$

2. Suppose  $z_1$  is an indicator representing whether or not an individual lives in an area in which there was a recent policy change to promote the importance of education. Could  $z_1$  be correlated with other unobservables captured in the error term?

Yes. Living in an area with recent education policy changes could correlate to a number of variables that relate to the earnings in that area, such as a person's trust in educational and governmental institutions, or the willingness of people to pay taxes to support education spending. As with most instrumental variables, the assumption that it is truly uncorrelated with the error term can always be questioned, although it may represent an improvement in estimating coefficients over straightforward OLS.

3. Using the same specification as that in [Question 4.5](#), estimate the equation by 2SLS, using both  $z_1$  and  $z_2$  as instrument variables. Interpret the results. How does the coefficient estimate on education change?

```
first_stage_a <- lm(education ~ z1, data = data )
first_stage_b <- lm(education ~ z2, data = data)
second_stage_a <- lm(data$logWage ~ first_stage_a$fitted + data$experience +
                      data$experienceSquare + data$raceColor +
                      data$dad_education + data$mom_education + data$rural +
                      data$city)
second_stage_b <- lm(data$logWage ~ first_stage_b$fitted + data$experience +
                      data$experienceSquare + data$raceColor +
                      data$dad_education + data$mom_education + data$rural +
                      data$city)
```

Table 14: Step One Regression Summary

	<i>Dependent variable:</i>	
	education	
	(1)	(2)
IV 1	0.242 (0.174)	
IV 2		1.006*** (0.187)
Constant	13.113*** (0.114)	12.529*** (0.158)
F Statistic	1.926	28.877***
df	1; 998	1; 998
Observations	1,000	1,000
R <sup>2</sup>	0.002	0.029
Adjusted R <sup>2</sup>	0.001	0.028
Residual Std. Error	2.728	2.690

·p<0.1; \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

Note that  $z_1$  is not correlated with education, and is thus not an appropriate instrumental variable to estimate the effect of education on wages.

Table 15: Step Two Regression summary

<i>Dependent variable:</i>	
	log(Wages)
Education ( $z_2$ )	0.045 (0.033)
Experience	0.072*** (0.013)
Expereince <sup>2</sup>	-0.003*** (0.001)
Race (White or Non-white)	-0.245*** (0.045)
Father's Education	0.006 (0.006)
Mother's Education	0.022** (0.007)
Rural (Yes or No)	-0.096** (0.034)
City (Yes or No)	0.191*** (0.034)
Constant	4.974*** (0.434)
F Statistic	21.634***
df	8; 714
Observations	723
R <sup>2</sup>	0.198
Adjusted R <sup>2</sup>	0.189
Residual Std. Error	0.398

·p<0.1; \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

The results of the regression show that the magnitude of the covariate coefficients is similar using the instrumental variable approach. The coefficient for education has dramatically decreased and is no longer statistically significant.

## Question 5: CLM 2

The dataset, `wealthy_candidates.csv`, contains candidate level electoral data from a developing country. Politically, each region (which is a subset of the country) is divided into smaller electoral districts where the candidate with the most votes wins the seat. This dataset has data on the financial wealth and electoral performance (voteshare) of electoral candidates. We are interested in understanding whether or not wealth is an electoral advantage. In other words, do wealthy candidates fare better in elections than their less wealthy peers?

1. Begin with a parsimonious, yet appropriate, specification. Why did you choose this model? Are your results statistically significant? Based on these results, how would you answer the research question? Is there a linear relationship between wealth and electoral performance?

In order to decide on a parsimonious model and discover any potential issues with data quality or variable distributions, we will first explore summary statistics and conduct a univariate analysis.

Table 16: Summary Statistics for Voting Data

Statistic	X	urb	lit	voteshare	absolute_wealth	logWealth
N	2,497	2,497	2,497	2,497	2,497	2,497
Mean	1,249.930	0.187	0.451	0.288	5,034,105.000	11.961
St. Dev.	721.080	0.149	0.092	0.123	31,098,493.000	5.388
Min	1	0.028	0.242	0.006	2.000	0.693
Pctl(25)	626	0.084	0.385	0.200	187,500.000	12.142
Median	1,250	0.147	0.460	0.293	1,336,629.000	14.106
Pctl(75)	1,874	0.243	0.510	0.368	4,092,001.000	15.225
Max	2,498	0.802	0.652	0.693	1,216,399,232.000	20.919

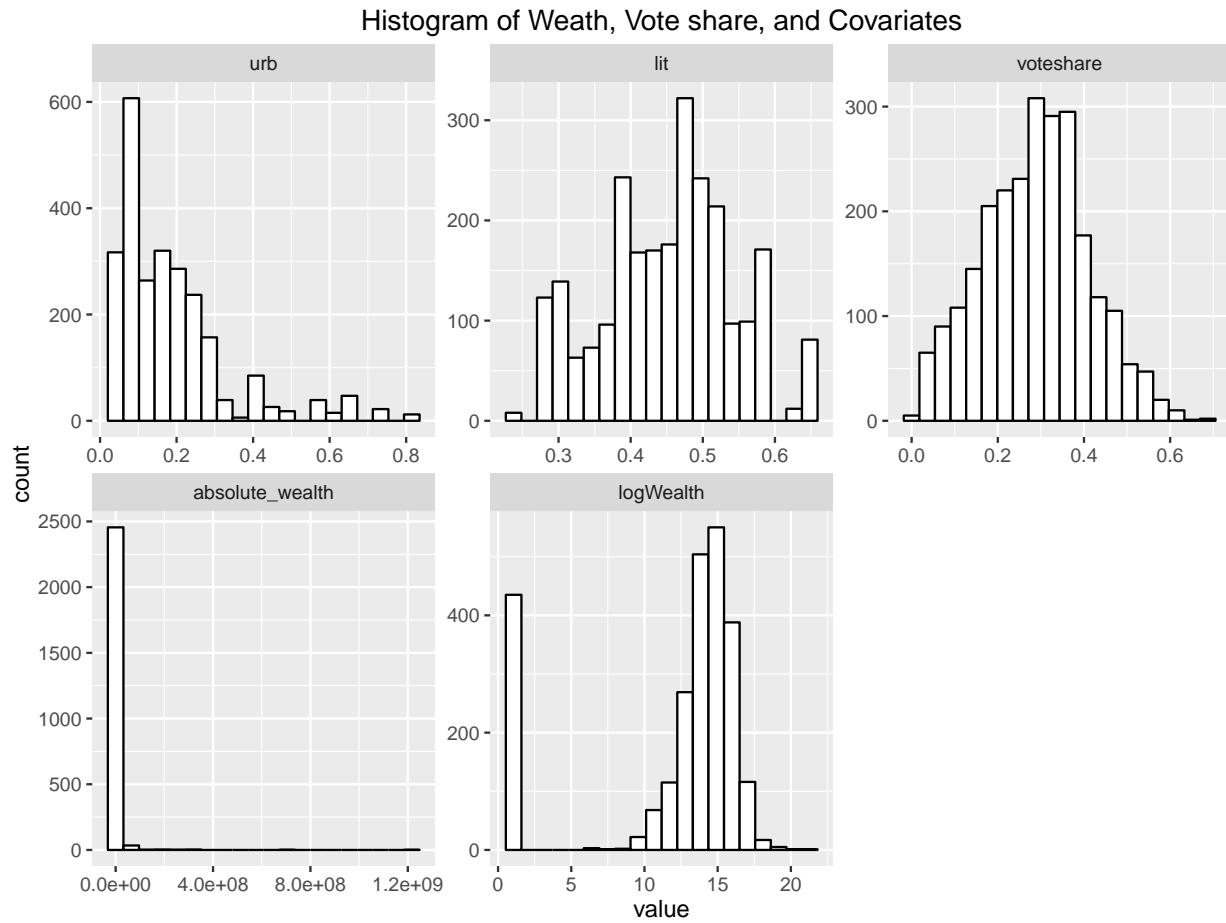


Figure 13: Histogram of Urban, Literacy, Voteshar, and Wealth Variables

Examining the absolute\_wealth and log(absolute\_wealth) variables reveals that a large number of observations have the same value. It seems likely that this value codes for having zero absolute wealth. Because we are interested in the effect of wealth on voteshare, it seems reasonable to look at this effect among those with wealth greater than zero. By creating a factor variable representing wealth greater than zero, the equation for our parsimonious model will be

$$\text{voteshare} = \beta_0 + \beta_1 \log(\text{wealth}) + \beta_2 \text{hasWealth}$$



Figure 14: Bar Plot of Wealth and No Wealth

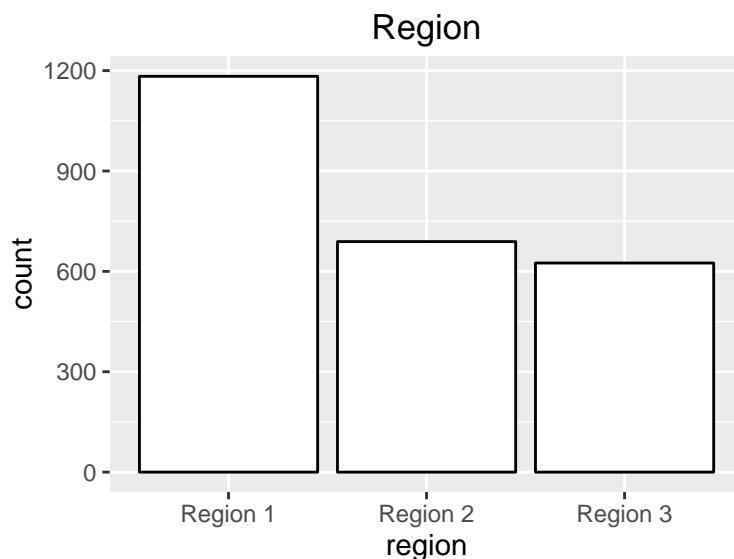


Figure 15: Bar Plot of Region

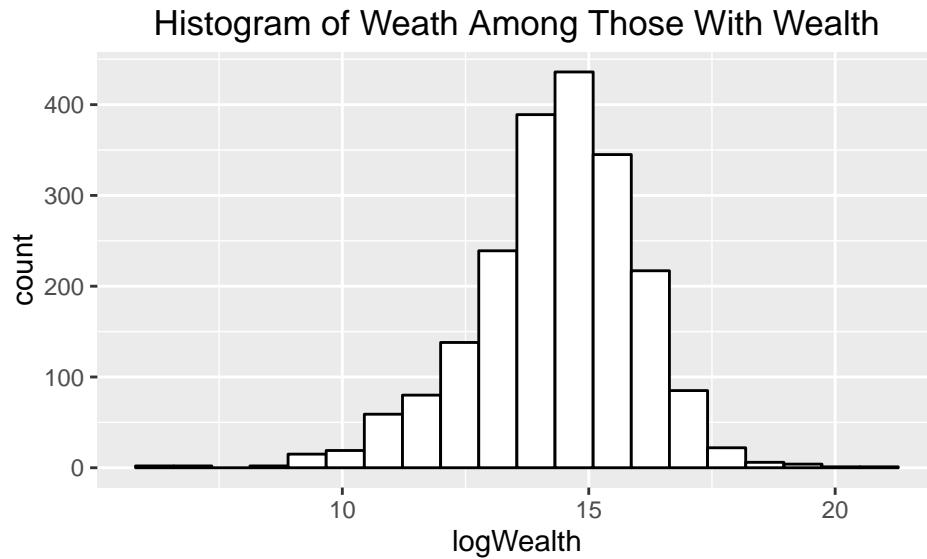


Figure 16: Histogram of log(Wealth) Among Those With Wealth

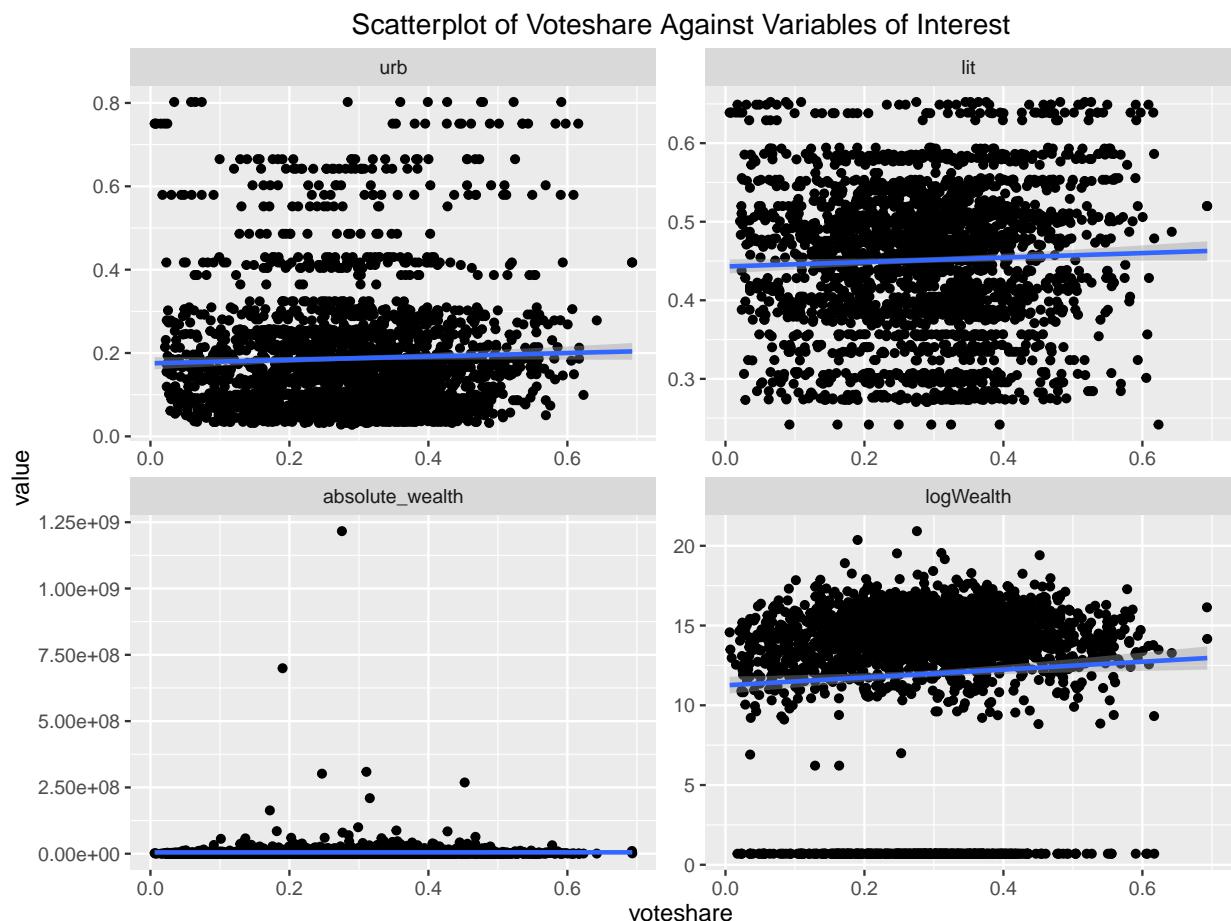


Figure 17: Scatterplots of Voteshare against Urban, Literacy, Absolute Wealth, and log(Absolute Wealth)

Table 17: Regression summary - Parsimonious Model

<i>Dependent variable:</i>	
	Voteshare
log(Wealth)	0.005** (0.002)
Has Wealth = Yes	-0.057* (0.026)
Constant	0.273*** (0.006)
F Statistic	7.189***
df	2; 2494
Observations	2,497
R <sup>2</sup>	0.006
Adjusted R <sup>2</sup>	0.005
Residual Std. Error	0.123

·p<0.1; \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

The parsimonious model reveals a very weak positive effect of wealth on voteshare. The effect is not of practical significance and suggests that there may not be a linear relationship between wealth and voteshare.

2. A team-member suggests adding a quadratic term to your regression. Based on your prior model, is such an addition warranted? Add this term and interpret the results. Do wealthier candidates fare better in elections?

Theoretically, the addition of a quadratic seems questionable because it suggests there is a point at which the returns to wealth stop increasing. Even if the quadratic model resulted in an improvement in the R<sup>2</sup> value, it may be difficult to justify the inclusion of a quadratic term without a theoretical reason.

Table 18: Regression summary - Quadratic Model

<i>Dependent variable:</i>	
	Voteshare
log(Wealth)	0.005** (0.002)
log(Wealth) <sup>2</sup>	-0.057* (0.026)
Has Wealth = Yes	0.273*** (0.006)
F Statistic	7.189***
df	2; 2494
Observations	2,497
R <sup>2</sup>	0.006
Adjusted R <sup>2</sup>	0.005
Residual Std. Error	0.123

·p<0.1; \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

3. Another team member suggests that it is important to take into account the fact that different regions have different electoral contexts. In particular, the relationship between candidate wealth and electoral performance might be different across states. Modify your model and report your results. Test the hypothesis that this addition is not needed.

First, we will look at how the variables of interest are distributed across regions.

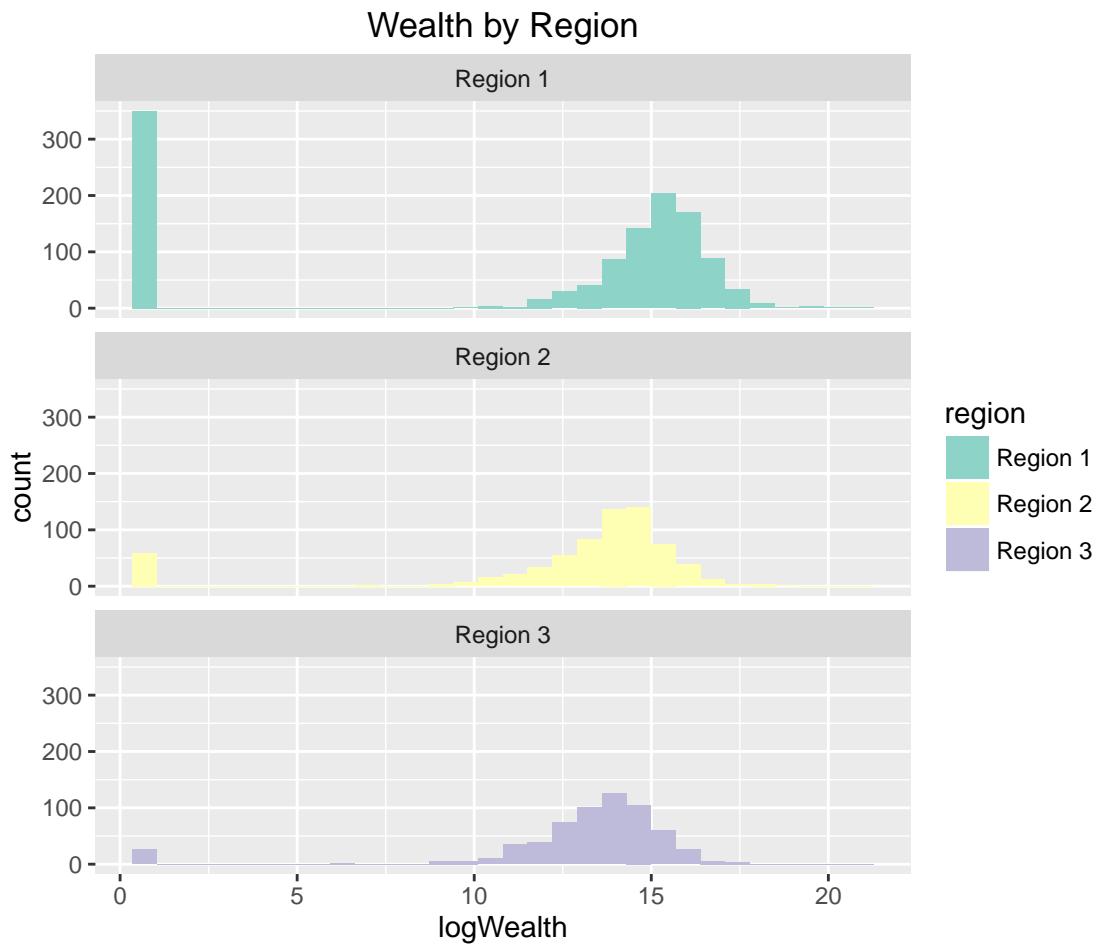


Figure 18: Histograms of Wealth By Region

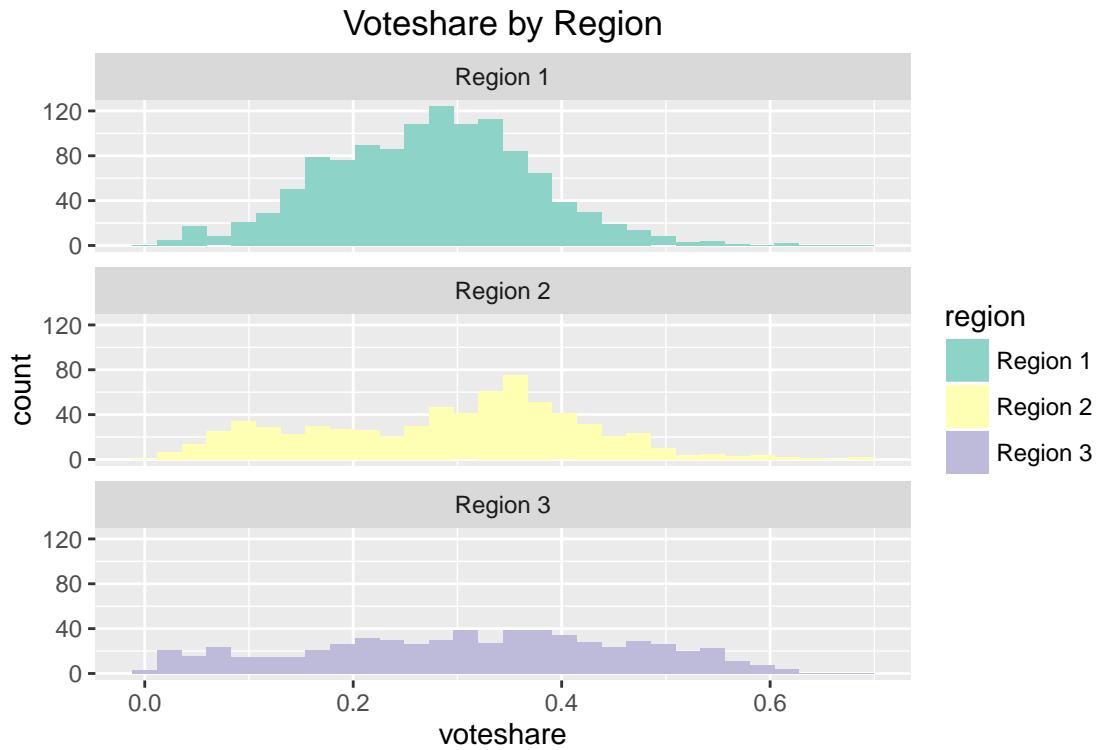


Figure 19: Histograms of Voteshare By Region

The histograms of wealth by region and voteshare by region do seem to reveal some genuine structural differences between the regions. The first region has much greater wealth inequality, with a large share of the observations having no wealth and greater average wealth among the wealthy than regions 2 and 3. The voteshare histogram for region one also suggests a difference in election structure, as voteshare percentages tend to be clustered between 15% and 40%, while regions two and three have relatively uniform voteshare distributions between 5% and 60%.

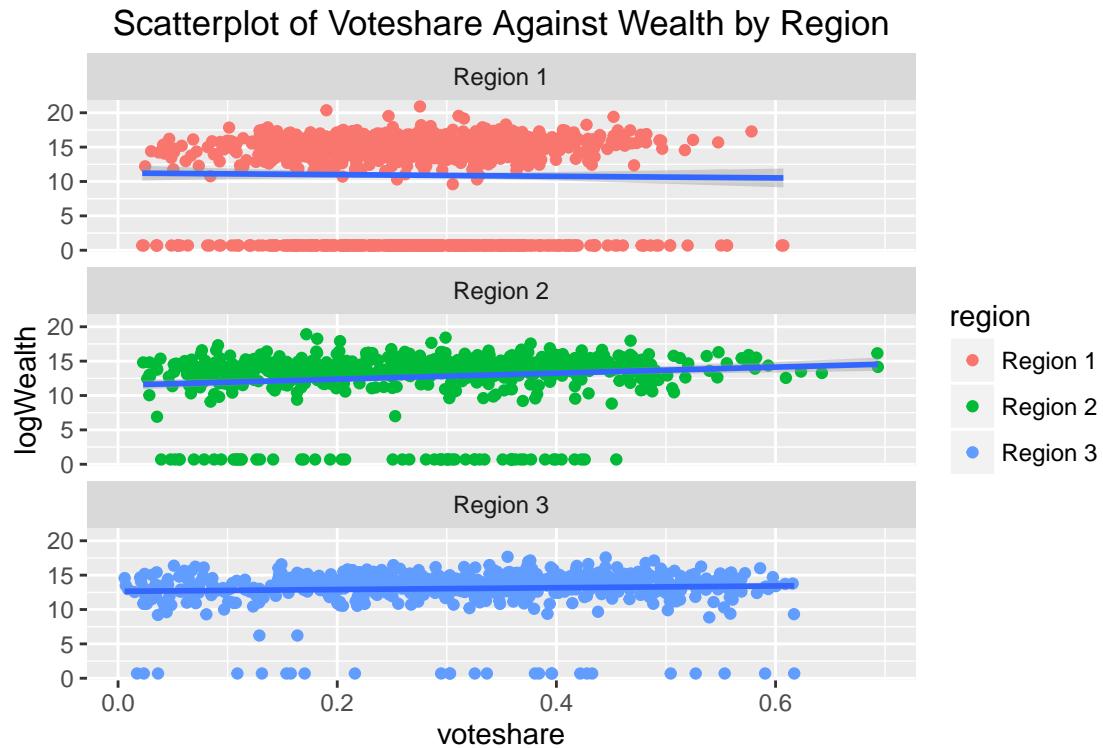


Figure 20: Scatterplots of Voteshare Against Wealth By Region

The scatter plots of voteshare against the wealth variables reveal that in all regions, there does not seem to be a strong linear relationship between wealth and voteshare. There seems to be a level of  $\log(\text{wealth})$  around 10, below which there are few observations, but once that floor is cleared, there doesn't seem to be a strong relationship between additional wealth and additional voteshare.

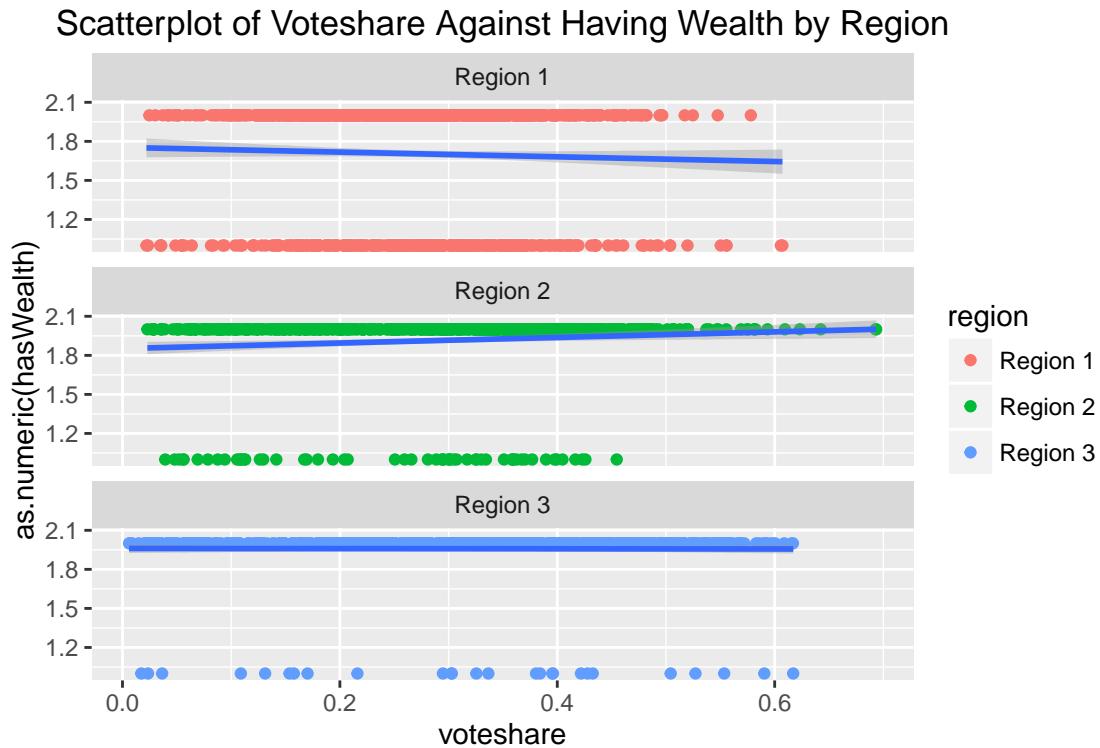


Figure 21: Scatterplots of Voteshare Against Having Wealth By Region

Table 19: Regression summary - Region Model

<i>Dependent variable:</i>	
	Voteshare
log(Wealth)	0.011*** (0.002)
Has Wealth = Yes	-0.156*** (0.029)
Region = 2	0.031*** (0.006)
Region = 3	0.055*** (0.007)
Constant	0.262*** (0.006)
F Statistic	19.367***
df	4; 2492
Observations	2,497
R <sup>2</sup>	0.032
Adjusted R <sup>2</sup>	0.030
Residual Std. Error	0.121

·p<0.1; \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

The regression output shows that the model including regions accounts for a greater share of the overall variation than the parsimonious model, with the difference between the two models being practically and

Table 20: Model Comparison

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
Parsimonious Model	2, 494	37.736				
Region Model	2, 492	36.741	2	0.995	33.755	0

statistically significant ( $F(2, 2492) = 33.755, p < 0.001$ ). This supports the notion that the structural differences in voteshare and wealth between the regions are important to take into account when estimating the effect of wealth on voteshare.

4. **Return to your parsimonious model. Do you think you have found a causal and unbiased estimate? Please state the conditions under which you would have an unbiased and causal estimates. Do these conditions hold?**

In order for OLS to make unbiased estimates of the OLS parameters, 4 conditions must hold:

1. The parameters have a linear relationship
2. Observations are observed through random sampling
3. There is no perfect collinearity among the independent variables
4. There is zero conditional mean of the error term and zero correlation of the error with any of the independent variables

And in order for our coefficients to represent causal effects, changes in the independent variable must not change the expected value of the error term. Ie.

$$\frac{\partial u}{\partial x} = 0$$

Two primary issues that lead me to conclude that we have not found a causal and unbiased estimate. The first is the lack of apparent linearity in the relationship between voteshare and wealth. This is supported by the low explanatory power of our model, as well as the scatter plot of wealth and voteshare that seems to imply non linear factors may be more important than linear. The second issue is that wealth is probably not exogenous, but are in fact correlated with a host of measures that are contained in the error term. For example, having more money may be correlated with being a well-known public figure, which may influence voteshare.

5. **Someone proposes a difference in difference design. Please write the equation for such a model. Under what circumstances would this design yield a causal effect?**

Supposing we had data collected at two different time periods, with  $\lambda_t$  being an indicator for the fixed time effect,  $a_i$  representing the time-fixed effect of being individual i, and  $u_{it}$  representing the time-variable portion of the error term. Then our population model would look like:

$$voteshare_{it} = \beta_0 + \lambda_t + \beta_1 \log(absolute\_wealth)_{it} + \beta_2 hasWealth_{it} + a_i + u_{it}$$

The first difference model would be obtained by subtracting the model at  $t_2$  from  $t_1$  and assuming parallel trends, ie that trends would have continued in the same way at  $t_2$  as in  $t_1$  in the absence of treatment.

$$\Delta voteshare_i = \lambda_t + \beta_1 \Delta \log(absolute\_wealth)_i + \beta_2 \Delta hasWealth_i + \Delta u_i$$

In order for the first difference model to be causal, changes in the difference in the independent variables must not change the expectation of changes in the error term.

$$\frac{\partial \Delta u}{\partial \Delta x} = 0$$

---

## Question 6: CLM 3

Your analytics team has been tasked with analyzing aggregate revenue, cost and sales data, which have been provided to you in the R workspace/data frame `retailSales.Rdata`.

Your task is two fold. First, your team is to develop a model for predicting (forecasting) revenues. Part of the model development documentation is a backtesting exercise where you train your model using data from the first two years and evaluate the model's forecasts using the last two years of data.

Second, management is equally interested in understanding variables that might affect revenues in support of management adjustments to operations and revenue forecasts. You are also to identify factors that affect revenues, and discuss how useful management's planned revenue is for forecasting revenues.

Your analysis should address the following:

- Exploratory Data Analysis: focus on bivariate and multivariate relationships.
- Be sure to assess conditions and identify unusual observations.

First we explore the whole dataset.

```
load("retailSales.Rdata")
data <- retailSales; rm(retailSales)
summary(data)

##      Year                  Product.line
##  Min.   :2004   Camping Equipment    :24108
##  1st Qu.:2005   Golf Equipment     : 8820
##  Median :2006   Mountaineering Equipment:12348
##  Mean   :2006   Outdoor Protection   : 8820
##  3rd Qu.:2006   Personal Accessories:30576
##  Max.   :2007

##
##      Product.type          Product
##  Eyewear       : 9408   Aloe Relief     :  588
##  Watches        : 7644   Astro Pilot     :  588
##  Lanterns       : 7056   Auto Pilot     :  588
##  Cooking Gear    : 5880   Bear Edge      :  588
##  Navigation       : 5880   Bear Survival Edge:  588
##  Climbing Accessories: 4116   Bella         :  588
##  (Other)          :44688   (Other)        :81144
##      Order.method.type  Retailer.country   Revenue
##  E-mail        :12096   Australia: 4032   Min.   :      0
##  Fax           :12096   Austria  : 4032   1st Qu.: 18579
##  Mail          :12096   Belgium   : 4032   Median : 59867
##  Sales visit    :12096   Brazil    : 4032   Mean   : 189418
##  Special        :12096   Canada    : 4032   3rd Qu.: 190193
##  Telephone       :12096   China     : 4032   Max.   :10054289
##  Web            :12096   (Other)   :60480   NA's    :59929
##      Planned.revenue   Product.cost      Quantity      Unit.cost
##  Min.   : 16   Min.   :     6   Min.   :    1   Min.   :  0.85
##  1st Qu.:19557 1st Qu.: 9432   1st Qu.: 328   1st Qu.: 11.43
```

```

## Median : 63907   Median : 32784   Median : 1043   Median : 36.83
## Mean   : 198818   Mean   : 111625   Mean   : 3607   Mean   : 84.89
## 3rd Qu.: 203996   3rd Qu.: 111371   3rd Qu.: 3288   3rd Qu.: 80.00
## Max.   :10054289   Max.   :6756853   Max.   :313628   Max.   :690.00
## NA's   :59929     NA's   :59929     NA's   :59929     NA's   :59929
##      Unit.price    Gross.profit    Unit.sale.price
## Min.   : 2.06     Min.   :-18160     Min.   : 0.00
## 1st Qu.: 23.00    1st Qu.: 8333     1st Qu.: 20.15
## Median : 66.77    Median : 25794    Median : 62.65
## Mean   : 155.99    Mean   : 77793    Mean   : 147.23
## 3rd Qu.: 148.30    3rd Qu.: 78254    3rd Qu.: 140.96
## Max.   :1359.72    Max.   :3521098   Max.   :1307.80
## NA's   :59929     NA's   :59929     NA's   :59929

```

The dataset contains 84,672 observations of 14 variables. 5 of them are categorical (`Product.line`, `Product.type`, `Product`, `Order.method.type`, `Retailer.country`), and `Year` should also be considered as categorical, since there are data from only 4 years (from 2004 to 2007).

```

data <- data %>%
  mutate(Year = as.factor(Year))

```

We also notice (from the output of `summary`) that some of the variables (all of them numerical) has a high number of NAs, the same in all cases (59929, i.e., 70.78% of the total number of observations). Do the NAs appear in the same observations for all those variables? Yes, they do.

```

# data_isNA <- as.data.frame(sapply(data, is.na))
data_isNA <- data %>% mutate_each(funs(is.na(.)))
head(data_isNA)

```

```

##      Year Product.line Product.type Product Order.method.type
## 1 FALSE      FALSE      FALSE    FALSE        FALSE
## 2 FALSE      FALSE      FALSE    FALSE        FALSE
## 3 FALSE      FALSE      FALSE    FALSE        FALSE
## 4 FALSE      FALSE      FALSE    FALSE        FALSE
## 5 FALSE      FALSE      FALSE    FALSE        FALSE
## 6 FALSE      FALSE      FALSE    FALSE        FALSE
##      Retailer.country Revenue Planned.revenue Product.cost Quantity Unit.cost
## 1          FALSE     FALSE      FALSE    FALSE    FALSE    FALSE
## 2          FALSE     FALSE      FALSE    FALSE    FALSE    FALSE
## 3          FALSE     TRUE      TRUE     TRUE    TRUE    TRUE
## 4          FALSE     TRUE      TRUE     TRUE    TRUE    TRUE
## 5          FALSE     FALSE      FALSE    FALSE    FALSE    FALSE
## 6          FALSE     TRUE      TRUE     TRUE    TRUE    TRUE
##      Unit.price Gross.profit Unit.sale.price
## 1      FALSE     FALSE      FALSE
## 2      FALSE     FALSE      FALSE
## 3      TRUE      TRUE      TRUE
## 4      TRUE      TRUE      TRUE
## 5      FALSE     FALSE      FALSE
## 6      TRUE      TRUE      TRUE

```

```
# vars_with_NAs <- apply(data_isNA, 2, sum)
vars_with_NAs <- data_isNA %>% summarise_each(funs(sum))
(vars_with_NAs <- names(vars_with_NAs)[vars_with_NAs>0])

## [1] "Revenue"           "Planned.revenue" "Product.cost"      "Quantity"
## [5] "Unit.cost"         "Unit.price"       "Gross.profit"     "Unit.sale.price"

sapply(data_isNA[, vars_with_NAs[-1]], identical,
       as.vector(data_isNA[, vars_with_NAs[1]]))

## Planned.revenue   Product.cost      Quantity      Unit.cost
##          TRUE        TRUE        TRUE        TRUE
##      Unit.price   Gross.profit Unit.sale.price
##          TRUE        TRUE        TRUE
```

And the amount of NAs per category is roughly the same for all categorical values (or at least there are non-missing data for all categories; below we just show the percentage per category for three of the numerical variables).

```
data_categorical <- data %>%
  select(which(names(data) %in% names(data)[sapply(data, is.factor)])) %>%
  mutate_each(funs(as.character(.))) %>% mutate(Revenue = data$Revenue)
data_categorical %>%
  select(Revenue, Year) %>%
  group_by(Year) %>%
  summarise_each(funs(100*mean(is.na(.)))) %>%
  rename("% of NAs in numerical variables" = Revenue)

## Source: local data frame [4 x 2]
##
##   Year % of NAs in numerical variables
##   (chr)          (dbl)
## 1 2004          67.95163
## 2 2005          65.49981
## 3 2006          71.70257
## 4 2007          77.95729

data_categorical %>%
  select(Revenue, Product.line) %>%
  group_by(Product.line) %>%
  summarise_each(funs(100*mean(is.na(.)))) %>%
  rename("% of NAs in numerical variables" = Revenue)

## Source: local data frame [5 x 2]
##
##   Product.line % of NAs in numerical variables
##   (chr)          (dbl)
## 1 Camping Equipment          65.26049
## 2 Golf Equipment            68.67347
## 3 Mountaineering Equipment 76.13379
## 4 Outdoor Protection        66.62132
## 5 Personal Accessories      74.77106
```

```

## Source: local data frame [21 x 2]
##
##   Retailer.country % of NAs in numerical variables
##   (chr)          (dbl)
## 1 Australia      77.15774
## 2 Austria        72.44544
## 3 Belgium        75.99206
## 4 Brazil         81.49802
## 5 Canada         57.66369
## 6 China          77.33135
## 7 Denmark        80.28274
## 8 Finland        79.46429
## 9 France         60.49107
## 10 Germany        59.37500
## 11 Italy          69.07242
## 12 Japan          58.60615
## 13 Korea          74.47917
## 14 Mexico         73.36310
## 15 Netherlands    70.03968
## 16 Singapore      70.70933
## 17 Spain          71.55258
## 18 Sweden          74.25595
## 19 Switzerland     80.03472
## 20 United Kingdom 70.23810
## 21 United States   52.28175

```

So we can omit all those missing observations (reducing our sample size to 24743), and continue with a further analysis of the numerical variables:

```

data <- data %>% na.omit()
data_categorical <- data %>%
  select(which(names(data) %in% names(data)[sapply(data, is.factor)]))
data_non_categorical <- data %>%
  select(which(names(data) %in% names(data)[!sapply(data, is.factor)]))
round(stat.desc(data_non_categorical, desc = TRUE, basic = TRUE), 2)

```

	Revenue	Planned.revenue	Product.cost	Quantity
## nbr.val	2.474300e+04	2.474300e+04	2.474300e+04	24743.00
## nbr.null	7.600000e+01	0.000000e+00	0.000000e+00	0.00
## nbr.na	0.000000e+00	0.000000e+00	0.000000e+00	0.00
## min	0.000000e+00	1.569000e+01	5.760000e+00	1.00
## max	1.005429e+07	1.005429e+07	6.756853e+06	313628.00
## range	1.005429e+07	1.005427e+07	6.756847e+06	313627.00
## sum	4.686776e+09	4.919342e+09	2.761941e+09	89237091.00
## median	5.986727e+04	6.390684e+04	3.278372e+04	1043.00
## mean	1.894182e+05	1.988175e+05	1.116251e+05	3606.56
## SE.mean	2.484130e+03	2.559050e+03	1.515680e+03	55.80
## CI.mean.0.95	4.869040e+03	5.015880e+03	2.970830e+03	109.38
## var	1.526863e+11	1.620349e+11	5.684198e+10	77048387.56
## std.dev	3.907509e+05	4.025355e+05	2.384156e+05	8777.72
## coef.var	2.060000e+00	2.020000e+00	2.140000e+00	2.43
##	Unit.cost	Unit.price	Gross.profit	Unit.sale.price
## nbr.val	24743.00	24743.00	2.474300e+04	24743.00

## nbr.null	0.00	0.00	0.000000e+00	76.00
## nbr.na	0.00	0.00	0.000000e+00	0.00
## min	0.85	2.06	-1.815960e+04	0.00
## max	690.00	1359.72	3.521098e+06	1307.80
## range	689.15	1357.66	3.539257e+06	1307.80
## sum	2100344.99	3859701.42	1.924835e+09	3642909.71
## median	36.83	66.77	2.579376e+04	62.65
## mean	84.89	155.99	7.779311e+04	147.23
## SE.mean	0.83	1.57	1.005230e+03	1.48
## CI.mean.0.95	1.63	3.08	1.970320e+03	2.89
## var	17190.71	60912.60	2.500267e+10	53846.65
## std.dev	131.11	246.80	1.581223e+05	232.05
## coef.var	1.54	1.58	2.030000e+00	1.58

All numerical variables are right-skewed, with long right tails (i.e., with several observations more than 2 standard deviations far from the mean), especially the ones corresponding to aggregate—non-unit—results.

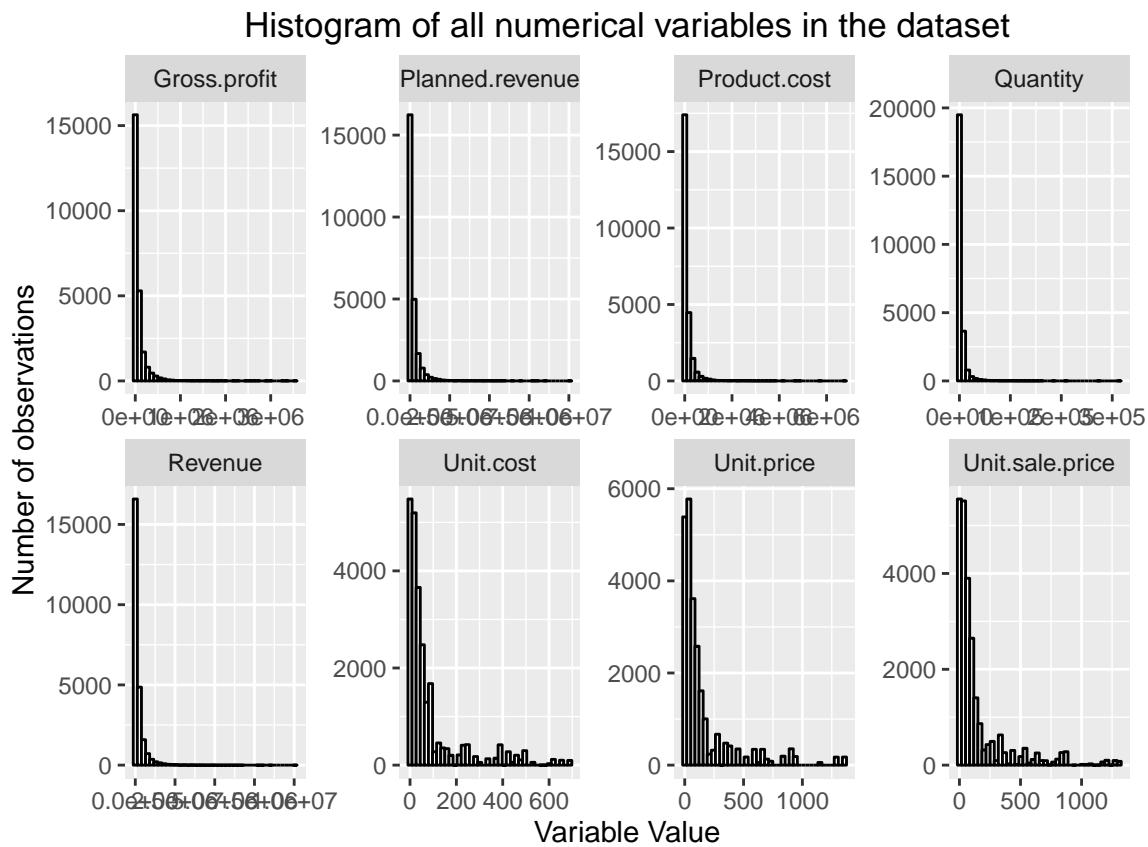


Figure 22: Histogram of all non-categorical variables in the dataset

Below we show the correlation matrix of the numerical variables, as well as two different representations of the scatterplot matrix (where we've used a sample of the data of size 500 because the plotting functions consume a lot of resources; that's why the correlations shown in the second Figure, only approximate, differ from the ones shown right below). As we might have expected, the correlations between `Revenue`,

`Planned.revenue`, `Product.cost`, and `Gross.profit` (i.e., the aggregate values), as well as those between `Unit.cost`, `Unit.price`, and `Unit.sale.price` (i.e., the values per unit), are positive and very high. `Quantity` is negatively correlated with the unitary variables (but that correlation is negligible in absolute value), and is moderately correlated ( $\rho \simeq 0.5$ ) with the aggregate values.

```
cor(data_non_categorical)
```

	Revenue	Planned.revenue	Product.cost	Quantity
## Revenue	1.000000	0.9990586	0.9903575	0.5055979
## Planned.revenue	0.9990586	1.0000000	0.9895792	0.4994770
## Product.cost	0.9903575	0.9895792	1.0000000	0.5061298
## Quantity	0.5055979	0.4994770	0.5061298	1.0000000
## Unit.cost	0.2463441	0.2550054	0.2415089	-0.1687497
## Unit.price	0.2332806	0.2421026	0.2194407	-0.1677662
## Gross.profit	0.9779407	0.9767878	0.9395732	0.4862920
## Unit.sale.price	0.2360448	0.2444078	0.2220105	-0.1674531
	Unit.cost	Unit.price	Gross.profit	Unit.sale.price
## Revenue	0.2463441	0.2332806	0.9779407	0.2360448
## Planned.revenue	0.2550054	0.2421026	0.9767878	0.2444078
## Product.cost	0.2415089	0.2194407	0.9395732	0.2220105
## Quantity	-0.1687497	-0.1677662	0.4862920	-0.1674531
## Unit.cost	1.0000000	0.9886870	0.2446187	0.9889263
## Unit.price	0.9886870	1.0000000	0.2456107	0.9992750
## Gross.profit	0.2446187	0.2456107	1.0000000	0.2485667
## Unit.sale.price	0.9889263	0.9992750	0.2485667	1.0000000

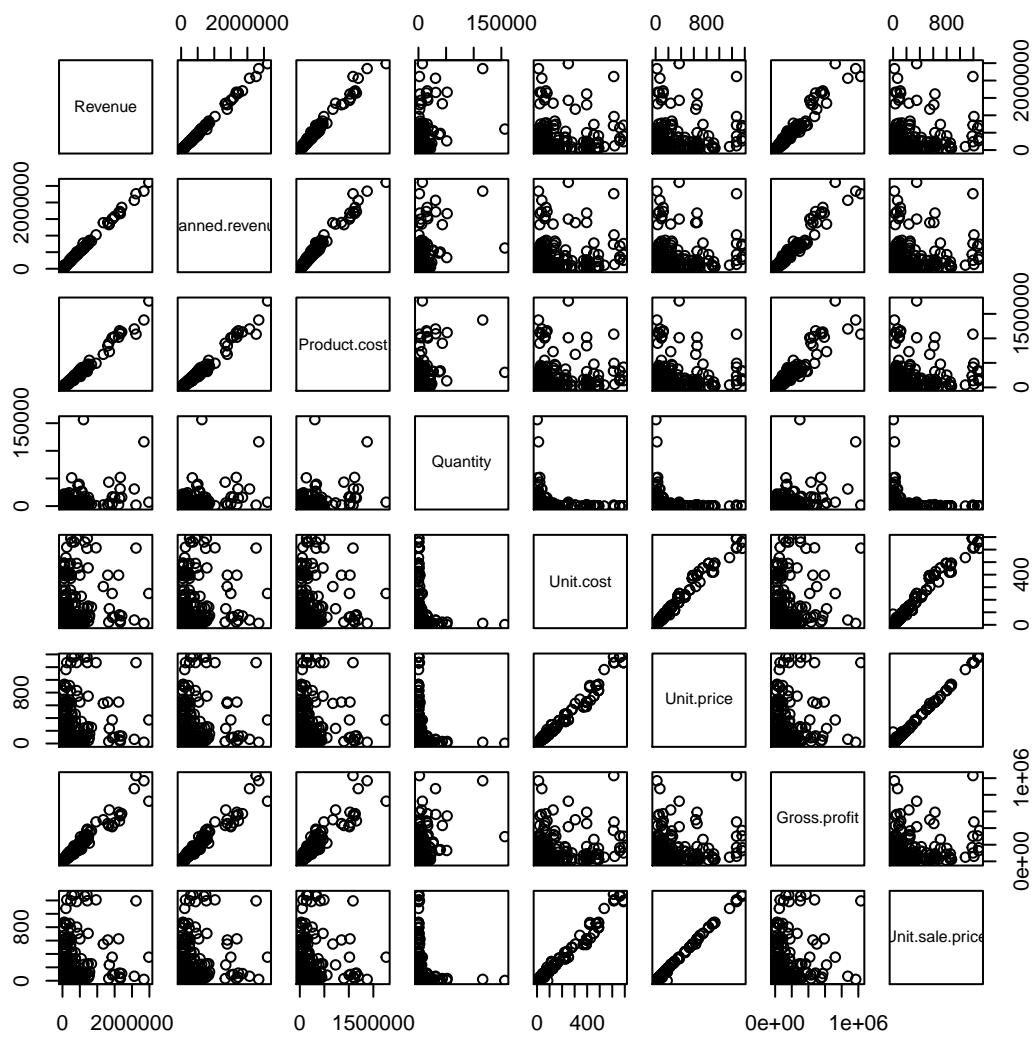


Figure 23: Scatterplot matrix of a sample of the dataset

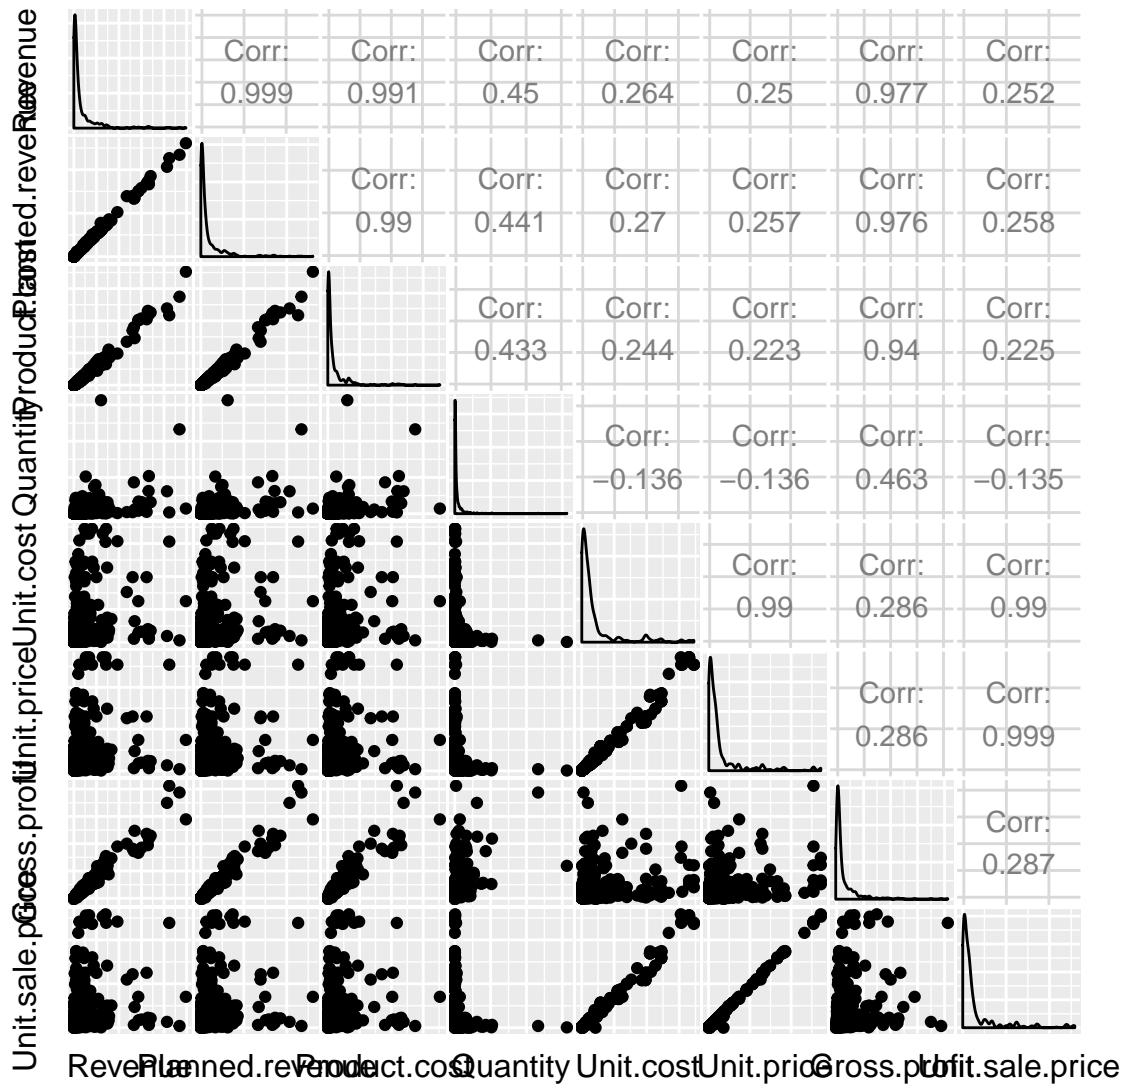


Figure 24: Scatterplot matrix of a sample of the dataset (with correlations)

We have to divide the dataset into two separate ones, to train and evaluate the model. Now we can convert `Year` back to a numerical variable (subtracting 2004 so the baseline is 0; that will make the intercept more intuitive when including `Year` in the regression model).

```
# Year back to integer (factor only useful for vizzes)
data <- data %>% mutate(Year2 = as.numeric(levels(Year))[Year] - 2004)
# One dataset per couple of years
# data_200405 <- data %>% filter(Year <= 2005)
# data_200607 <- data %>% filter(Year > 2005)
data_200405 <- data %>% filter(Year2 <= 1)
data_200607 <- data %>% filter(Year2 > 1)
```

Not all products appear in both periods so some re-factoring is needed, otherwise the `predict.lm` function will not work when it encounters a category not present before:

```
# Re-factor Product (since the levels differ by period)
products_200405 <- data.frame(Product = levels(droplevels(data_200405$Product)))
products_200607 <- data.frame(Product = levels(droplevels(data_200607$Product)))
continuing_products <- intersect(products_200405, products_200607)
(new_or_discontinued_products <- union(products_200405, products_200607) %>%
  setdiff(continuing_products))
```

```
##      Product
## 1 Trail Master
## 2 Trail Star
## 3 Auto Pilot

# Products present in one period and not the other are labelled as "Other"
data_200405 <- data_200405 %>%
  mutate(Product = ifelse(Product %in% new_or_discontinued_products$Product,
                         "Other", as.character(Product))) %>%
  mutate(Product = factor(Product))
data_200607 <- data_200607 %>%
  mutate(Product = ifelse(Product %in% new_or_discontinued_products$Product,
                         "Other", as.character(Product))) %>%
  mutate(Product = factor(Product))
```

There are some variables that are calculated from `Revenue` (or vice versa) so including them in any regression model would lead to a perfect fit. In particular, `Gross.profit = Revenue - Product.cost`. And `Revenue` should be equal to `Unit.sale.price` times `Quantity`, though this is not always the case, and there are differences in many cases (53.4% of the total number of observations).

```
head(data %>% select(Revenue, Product.cost, Gross.profit) %>%
  mutate(Revenue2 = Product.cost + Gross.profit))
```

```
##      Revenue Product.cost Gross.profit Revenue2
## 1 315044.33     158371.76    156672.57 315044.33
## 2 13444.68      6298.80     7145.88  13444.68
## 3 181120.24     89413.06    91707.18 181120.24
## 4 69608.15      35326.25    34281.90  69608.15
## 5 30940.35      16370.97    14569.38 30940.35
## 6 74321.18      36531.63    37789.55 74321.18
```

```

all(round(data$Revenue - data$Product.cost, 2) == round(data$Gross.profit, 2))

## [1] TRUE

head(data %>% select(Revenue, Unit.sale.price, Quantity) %>%
    mutate(Revenue2 = Unit.sale.price * Quantity))

##      Revenue Unit.sale.price Quantity Revenue2
## 1 315044.33      5.195714    66385 344917.49
## 2 13444.68       6.190000    2172  13444.68
## 3 181120.24      5.488000    35696 195899.65
## 4 69608.15       5.040000    15205  76633.20
## 5 30940.35       3.950000    7833   30940.35
## 6 74321.18       5.585000   14328  80021.88

```

But not only that: **in order to make predictions with our models, we need to use the values of the independent variables for which we want our predictions to be made; it seems very unlikely that the company will know in advance any of the numerical variables, same as the revenues are unknown.** The only exception is `Planned.revenue`, which already is a forecast (so we will use it alone, with no other variable, to evaluate how useful it is forecasting). (There might be other exceptions, such as `Unit.cost` and `Unit.price`, which could already be set by the beginning of the year and not vary over the year, but just in case we won't use them.)

```

data %>% filter(Product == "TX") %>% group_by(Year) %>%
    summarize(mean.Unit.Cost.TX = mean(Unit.cost),
              sd.Unit.Cost.TX = sd(Unit.cost),
              mean.Unit.price.TX = mean(Unit.price),
              sd.Unit.price.TX = sd(Unit.price))

## Source: local data frame [4 x 5]
##
##      Year mean.Unit.Cost.TX sd.Unit.Cost.TX mean.Unit.price.TX
##      (fctr)          (dbl)          (dbl)          (dbl)
## 1 2004        102.9073     0.5964767      188.0542
## 2 2005        103.6643     1.2687343      189.0911
## 3 2006        104.6195     1.7755464      191.6700
## 4 2007        104.8819     3.8325546      193.0520
## Variables not shown: sd.Unit.price.TX (dbl).

```

Let's continue the EDA, now with the categorical variables (the ones that can be used for forecasting). We explored several combinations, but for the sake of brevity we'll include here only the most interesting ones, starting with `Year` and `Order.method.type`:

Table 21: Percentage of Revenue per order method type by year

	2004	2005	2006	2007
E-mail	10.4%	3.8%	1.6%	1.5%
Fax	3.1%	1.7%	0.9%	0.7%
Mail	2.5%	1.4%	0.5%	0.0%
Sales visit	11.1%	6.9%	4.9%	5.0%
Special	1.5%	0.9%	0.1%	0.1%
Telephone	19.6%	9.2%	2.5%	1.6%
Web	51.8%	76.0%	89.6%	91.1%

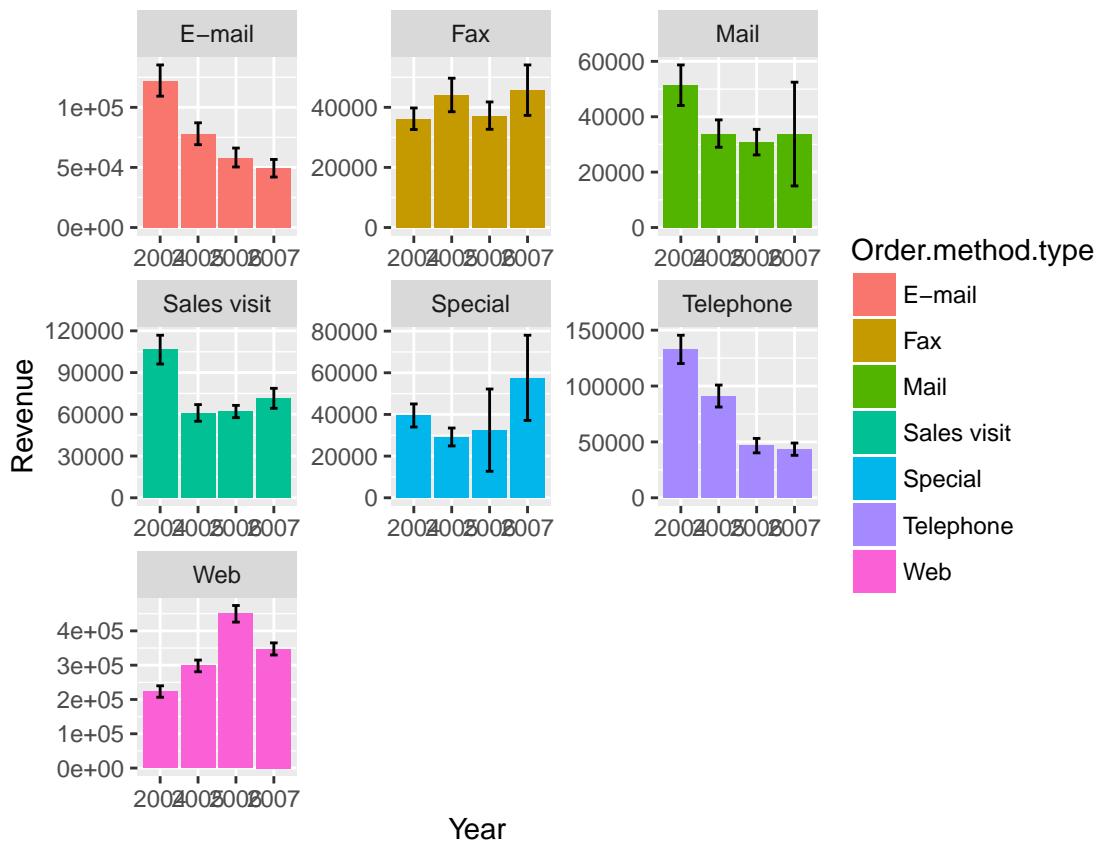


Figure 25: Bar chart of aggregate revenues per Order method type and Year

```

params = c("Year2", "Order.method.type", "Planned.revenue")
model1 <- lm(as.formula(paste("Revenue", paste(params, sep = "", collapse = " * "),
                           sep = " ~ ")), data_200405)
model1 <- lm(Revenue ~ Planned.revenue + Year2 * Order.method.type,
            data_200405)

```

```
# coefest(model1, vcov = vcovHC)
model1_predictions <- predict(model1, data_200607[, params],
                             interval = "prediction")

stargazer2(model1,
            title = paste("Summary of the regression model of revenue on year",
                          "and order method type"))
```

Table 22: Summary of the regression model of revenue on year and order method type

	<i>Dependent variable:</i>
	Revenue
Planned.revenue	0.970*** (0.003)
Year2	2,871.262*** (525.115)
Order.method.typeFax	4,038.016*** (534.831)
Order.method.typeMail	2,241.022*** (561.078)
Order.method.typeSales visit	−243.732 (535.810)
Order.method.typeSpecial	3,432.403*** (543.549)
Order.method.typeTelephone	−871.942 (625.014)
Order.method.typeWeb	2,892.816*** (605.795)
Year2:Order.method.typeFax	−2,875.152*** (581.253)
Year2:Order.method.typeMail	−1,233.207* (610.986)
Year2:Order.method.typeSales visit	461.429 (615.616)
Year2:Order.method.typeSpecial	−2,178.338*** (585.654)
Year2:Order.method.typeTelephone	−451.167 (743.353)
Year2:Order.method.typeWeb	−4,820.421*** (795.420)
Constant	−5,189.293*** (564.119)
F Statistic	51,978.079***
df	14; 14072
Observations	14,087
R <sup>2</sup>	0.998
Adjusted R <sup>2</sup>	0.998
Residual Std. Error	14,563.800

·p<0.1; \*p<0.05; \*\*p<0.01; \*\*\*p<0.001

We could use both the *Bayesian Information Criterion* (BIC) and the *Akaike Information Criterion* (AIC) to compare different models (preferring the one with the lowest BIC or AIC). Both metrics introduce a penalty term for the number of parameters in the model (larger in BIC), though they don't answer exactly the same question: the AIC tries to select the model that most adequately describes an unknown, high dimensional reality, i.e., it is aimed for situations where we don't necessarily think there's a model so much as a bunch of effects of different sizes, and we want to get good prediction errors. On the contrary, the BIC tries to find the "true" model among the set of candidates. For that reason, the AIC may be better for our objective.

But these metrics evaluate the models (based on the data used to build them), not their predictions: to evaluate their accuracy, we'll use the RMSE and the MAE (*Mean Absolute Error*; the MAPE cannot be used here because `Revenue = 0` in some cases). The greater the difference between them, the greater the variance in the prediction errors in the test set.

```
cbind(BIC = BIC(model1), AIC = AIC(model1))

##           BIC      AIC
## [1,] 310199.3 310078.4

rbind("model1" = accuracy(model1_predictions[, 1],
                           data_200607$Revenue)[, c("RMSE", "MAE")])

##           RMSE      MAE
## model1 19712.78 9370.763
```

Let's continue exploring `Retailer.country` and `Product.line`:

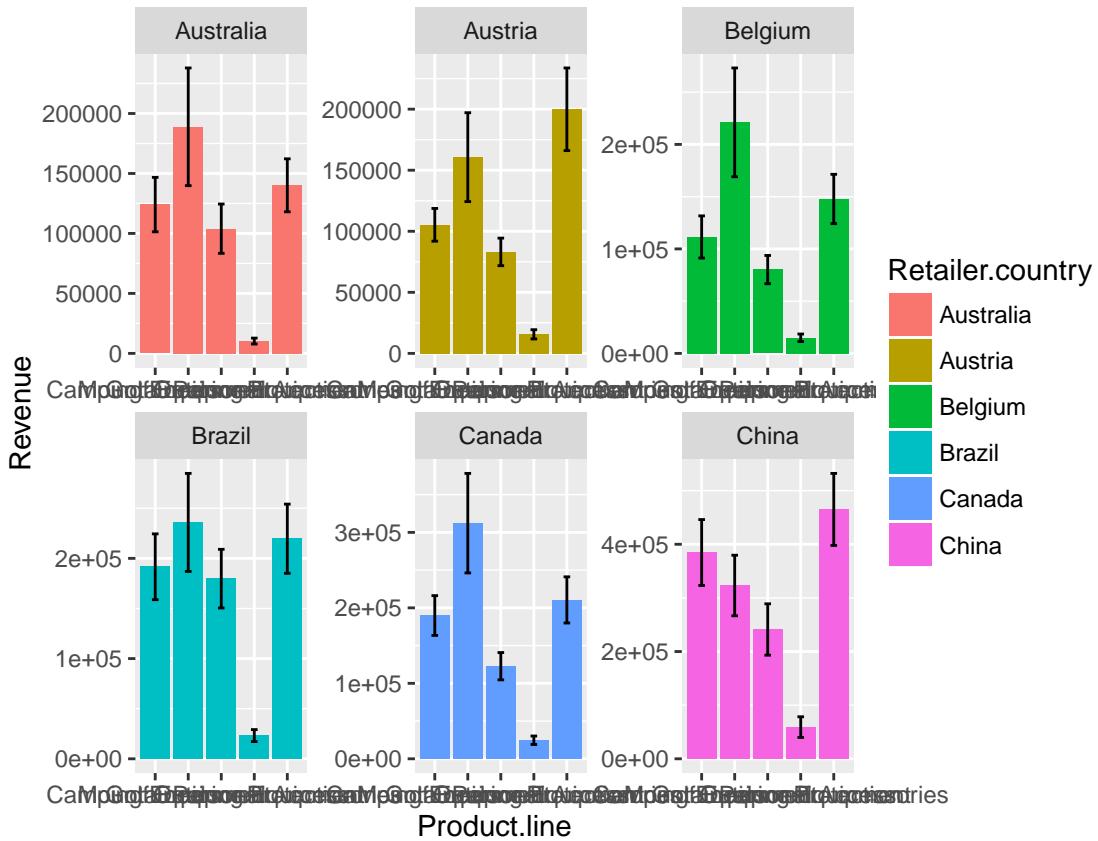


Figure 26: Bar chart of aggregate revenues per Retailer country (only 6 shown) and Product Line

The revenue per product line seems quite similar for all countries, so the interaction may be light: we'll explore both models anyway.

```
params = c("Year2", "Order.method.type", "Planned.revenue", "Retailer.country",
          "Product.line")
model2 <- lm(as.formula(paste("Revenue", paste(params[3], "+", params[1], "*",
                                         params[2], "+", params[4], "+",
                                         params[5]), sep = " ~ "))),
            data_200405)
coeftest(model2, vcov = vcovHC) [, c(1, 4)]
```

	Estimate	Pr(> t )
## (Intercept)	-6097.6927639	9.784900e-23
## Planned.revenue	0.9726947	0.000000e+00
## Year2	2208.6066847	3.275944e-05
## Order.method.typeFax	2457.9651227	4.930208e-06
## Order.method.typeMail	3432.4006307	3.007726e-09
## Order.method.typeSales visit	-202.2007356	7.034878e-01

## Order.method.typeSpecial	2080.0060929	3.054700e-04
## Order.method.typeTelephone	-2045.5292460	8.901631e-04
## Order.method.typeWeb	1088.9459239	5.770631e-02
## Retailer.countryAustria	1224.0005008	1.069868e-02
## Retailer.countryBelgium	206.0187233	6.677030e-01
## Retailer.countryBrazil	-658.0191830	2.845138e-01
## Retailer.countryCanada	-1878.2817175	7.035228e-04
## Retailer.countryChina	-5197.6686614	5.830496e-07
## Retailer.countryDenmark	1213.1740800	7.871759e-03
## Retailer.countryFinland	-2627.6633738	9.609300e-05
## Retailer.countryFrance	-1566.2696181	4.141165e-03
## Retailer.countryGermany	-2513.2700378	1.710364e-05
## Retailer.countryItaly	-976.1927286	5.451211e-02
## Retailer.countryJapan	-2939.4246355	6.135175e-06
## Retailer.countryKorea	-3497.6142328	2.710083e-08
## Retailer.countryMexico	-2045.7740039	7.536002e-05
## Retailer.countryNetherlands	-1136.2036667	1.682214e-02
## Retailer.countrySingapore	-2274.3549840	5.791529e-05
## Retailer.countrySpain	-1016.4012984	4.761992e-02
## Retailer.countrySweden	184.5112399	7.051869e-01
## Retailer.countrySwitzerland	353.8816929	4.690483e-01
## Retailer.countryUnited Kingdom	-3304.3008012	1.205611e-05
## Retailer.countryUnited States	-6511.4804448	6.872598e-16
## Product.lineGolf Equipment	-5623.0748648	4.852857e-19
## Product.lineMountaineering Equipment	4387.1172536	1.072891e-49
## Product.lineOutdoor Protection	6487.1447063	2.617534e-69
## Product.linePersonal Accessories	9336.8942369	1.544458e-257
## Year2:Order.method.typeFax	-1927.5772999	2.763264e-03
## Year2:Order.method.typeMail	-1083.5909982	1.002371e-01
## Year2:Order.method.typeSales visit	496.9065320	4.346150e-01
## Year2:Order.method.typeSpecial	-1093.5688051	1.198593e-01
## Year2:Order.method.typeTelephone	1228.4141391	1.055746e-01
## Year2:Order.method.typeWeb	-4657.8209711	2.031962e-09

```
model2_predictions <- predict(model2, data_200607[, params],
                               interval = "prediction")
```

```
cbind(BIC(model1, model2),
      AIC = AIC(model1, model2)[, 2])
```

```
##          df      BIC      AIC
## model1 16 310199.3 310078.4
## model2 40 308373.7 308071.6
```

```
rbind("model1" = accuracy(model1_predictions[, 1],
                           data_200607$Revenue)[, c("RMSE", "MAE")],
      "model2" = accuracy(model2_predictions[, 1],
                           data_200607$Revenue)[, c("RMSE", "MAE")])
```

```
##          RMSE      MAE
## model1 19712.78 9370.763
## model2 18265.22 9550.499
```

```

model3 <- lm(as.formula(paste("Revenue", paste(params[3], "+", params[1], "*",
                                              params[2], "+", params[4], "*",
                                              params[5]), sep = " ~ ")),
             data_200405)
coeftest(model3, vcov = vcovHC)[, c(1, 4)]
```

	Estimate
##	-4.126412e+03
## (Intercept)	9.728107e-01
## Planned.revenue	2.161037e+03
## Year2	2.467050e+03
## Order.method.typeFax	3.465538e+03
## Order.method.typeMail	-1.479667e+02
## Order.method.typeSales visit	2.284567e+03
## Order.method.typeSpecial	-2.026628e+03
## Order.method.typeTelephone	1.172547e+03
## Order.method.typeWeb	8.169051e+01
## Retailer.countryAustria	-2.844152e+02
## Retailer.countryBelgium	-2.823049e+03
## Retailer.countryBrazil	-3.775110e+03
## Retailer.countryCanada	-1.167779e+04
## Retailer.countryChina	1.360433e+03
## Retailer.countryDenmark	-4.063831e+03
## Retailer.countryFinland	-3.668372e+03
## Retailer.countryFrance	-4.471593e+03
## Retailer.countryGermany	-1.770205e+03
## Retailer.countryItaly	-5.392263e+03
## Retailer.countryJapan	-5.104382e+03
## Retailer.countryKorea	-3.050416e+03
## Retailer.countryMexico	-2.308742e+03
## Retailer.countryNetherlands	-3.600832e+03
## Retailer.countrySingapore	-2.200762e+03
## Retailer.countrySpain	2.080229e+02
## Retailer.countrySweden	2.513635e+01
## Retailer.countrySwitzerland	-7.226318e+03
## Retailer.countryUnited Kingdom	-1.334940e+04
## Retailer.countryUnited States	-1.924985e+03
## Product.lineGolf Equipment	3.420156e+02
## Product.lineMountaineering Equipment	2.766749e+03
## Product.lineOutdoor Protection	3.637943e+03
## Product.linePersonal Accessories	-1.659733e+03
## Year2:Order.method.typeFax	-8.444572e+02
## Year2:Order.method.typeMail	5.878369e+02
## Year2:Order.method.typeSales visit	-1.551288e+03
## Year2:Order.method.typeSpecial	1.239126e+03
## Year2:Order.method.typeTelephone	-4.573377e+03
## Year2:Order.method.typeWeb	-5.806299e+02
## Retailer.countryAustria:Product.lineGolf Equipment	-8.297121e+02
## Retailer.countryBelgium:Product.lineGolf Equipment	-1.611695e+03
## Retailer.countryBrazil:Product.lineGolf Equipment	-2.598464e+03
## Retailer.countryCanada:Product.lineGolf Equipment	-1.891069e+03
## Retailer.countryChina:Product.lineGolf Equipment	8.074872e+02
## Retailer.countryDenmark:Product.lineGolf Equipment	-1.080670e+04

## Retailer.countryFrance:Product.lineGolf Equipment	-2.391424e+03
## Retailer.countryGermany:Product.lineGolf Equipment	-6.077728e+03
## Retailer.countryItaly:Product.lineGolf Equipment	-4.040289e+03
## Retailer.countryJapan:Product.lineGolf Equipment	-4.073835e+03
## Retailer.countryKorea:Product.lineGolf Equipment	-4.021083e+03
## Retailer.countryMexico:Product.lineGolf Equipment	-2.037534e+03
## Retailer.countryNetherlands:Product.lineGolf Equipment	-2.310704e+03
## Retailer.countrySingapore:Product.lineGolf Equipment	-2.672779e+03
## Retailer.countrySpain:Product.lineGolf Equipment	-1.374177e+03
## Retailer.countrySweden:Product.lineGolf Equipment	-5.119653e+03
## Retailer.countrySwitzerland:Product.lineGolf Equipment	6.618909e+02
## Retailer.countryUnited Kingdom:Product.lineGolf Equipment	-5.538798e+03
## Retailer.countryUnited States:Product.lineGolf Equipment	-1.181372e+04
## Retailer.countryAustria:Product.lineMountaineering Equipment	1.882310e+03
## Retailer.countryBelgium:Product.lineMountaineering Equipment	9.824330e+02
## Retailer.countryBrazil:Product.lineMountaineering Equipment	3.701488e+03
## Retailer.countryCanada:Product.lineMountaineering Equipment	3.253141e+03
## Retailer.countryChina:Product.lineMountaineering Equipment	1.039267e+04
## Retailer.countryDenmark:Product.lineMountaineering Equipment	7.229459e+02
## Retailer.countryFinland:Product.lineMountaineering Equipment	3.770139e+03
## Retailer.countryFrance:Product.lineMountaineering Equipment	3.959380e+03
## Retailer.countryGermany:Product.lineMountaineering Equipment	4.200868e+03
## Retailer.countryItaly:Product.lineMountaineering Equipment	2.052448e+03
## Retailer.countryJapan:Product.lineMountaineering Equipment	1.996780e+03
## Retailer.countryKorea:Product.lineMountaineering Equipment	3.648041e+03
## Retailer.countryMexico:Product.lineMountaineering Equipment	2.868689e+03
## Retailer.countryNetherlands:Product.lineMountaineering Equipment	3.100629e+03
## Retailer.countrySingapore:Product.lineMountaineering Equipment	3.324492e+03
## Retailer.countrySpain:Product.lineMountaineering Equipment	1.571317e+03
## Retailer.countrySweden:Product.lineMountaineering Equipment	2.087557e+03
## Retailer.countrySwitzerland:Product.lineMountaineering Equipment	1.719986e+03
## Retailer.countryUnited Kingdom:Product.lineMountaineering Equipment	7.712514e+03
## Retailer.countryUnited States:Product.lineMountaineering Equipment	1.036049e+04
## Retailer.countryAustria:Product.lineOutdoor Protection	9.136473e+02
## Retailer.countryBelgium:Product.lineOutdoor Protection	6.263762e+02
## Retailer.countryBrazil:Product.lineOutdoor Protection	3.650149e+03
## Retailer.countryCanada:Product.lineOutdoor Protection	3.179778e+03
## Retailer.countryChina:Product.lineOutdoor Protection	9.386646e+03
## Retailer.countryDenmark:Product.lineOutdoor Protection	-3.303281e+02
## Retailer.countryFinland:Product.lineOutdoor Protection	4.024804e+03
## Retailer.countryFrance:Product.lineOutdoor Protection	3.303060e+03
## Retailer.countryGermany:Product.lineOutdoor Protection	3.558898e+03
## Retailer.countryItaly:Product.lineOutdoor Protection	1.911161e+03
## Retailer.countryJapan:Product.lineOutdoor Protection	4.840081e+03
## Retailer.countryKorea:Product.lineOutdoor Protection	4.407140e+03
## Retailer.countryMexico:Product.lineOutdoor Protection	3.070475e+03
## Retailer.countryNetherlands:Product.lineOutdoor Protection	2.085135e+03
## Retailer.countrySingapore:Product.lineOutdoor Protection	3.462718e+03
## Retailer.countrySpain:Product.lineOutdoor Protection	1.804383e+03
## Retailer.countrySweden:Product.lineOutdoor Protection	6.220915e+02
## Retailer.countrySwitzerland:Product.lineOutdoor Protection	9.446517e+00
## Retailer.countryUnited Kingdom:Product.lineOutdoor Protection	6.924935e+03
## Retailer.countryUnited States:Product.lineOutdoor Protection	1.152421e+04
## Retailer.countryAustria:Product.linePersonal Accessories	3.283102e+03

## Retailer.countryBelgium:Product.linePersonal Accessories	2.167725e+03
## Retailer.countryBrazil:Product.linePersonal Accessories	5.369215e+03
## Retailer.countryCanada:Product.linePersonal Accessories	5.289153e+03
## Retailer.countryChina:Product.linePersonal Accessories	1.511299e+04
## Retailer.countryDenmark:Product.linePersonal Accessories	5.129432e+02
## Retailer.countryFinland:Product.linePersonal Accessories	5.443902e+03
## Retailer.countryFrance:Product.linePersonal Accessories	5.518373e+03
## Retailer.countryGermany:Product.linePersonal Accessories	6.109236e+03
## Retailer.countryItaly:Product.linePersonal Accessories	3.107783e+03
## Retailer.countryJapan:Product.linePersonal Accessories	6.781199e+03
## Retailer.countryKorea:Product.linePersonal Accessories	4.233633e+03
## Retailer.countryMexico:Product.linePersonal Accessories	2.538880e+03
## Retailer.countryNetherlands:Product.linePersonal Accessories	3.174470e+03
## Retailer.countrySingapore:Product.linePersonal Accessories	3.281874e+03
## Retailer.countrySpain:Product.linePersonal Accessories	3.009094e+03
## Retailer.countrySweden:Product.linePersonal Accessories	1.788856e+03
## Retailer.countrySwitzerland:Product.linePersonal Accessories	1.039136e+03
## Retailer.countryUnited Kingdom:Product.linePersonal Accessories	9.577319e+03
## Retailer.countryUnited States:Product.linePersonal Accessories	1.794006e+04
## Pr(> t )	
## (Intercept)	5.872317e-11
## Planned.revenue	0.000000e+00
## Year2	9.966808e-05
## Order.method.typeFax	1.435505e-05
## Order.method.typeMail	3.290722e-08
## Order.method.typeSales visit	7.909606e-01
## Order.method.typeSpecial	1.900530e-04
## Order.method.typeTelephone	1.161782e-03
## Order.method.typeWeb	4.425478e-02
## Retailer.countryAustria	8.759709e-01
## Retailer.countryBelgium	6.372966e-01
## Retailer.countryBrazil	2.766056e-03
## Retailer.countryCanada	1.449057e-06
## Retailer.countryChina	8.102252e-11
## Retailer.countryDenmark	6.716526e-03
## Retailer.countryFinland	4.339669e-05
## Retailer.countryFrance	4.873629e-08
## Retailer.countryGermany	1.230638e-09
## Retailer.countryItaly	4.202337e-03
## Retailer.countryJapan	8.985533e-10
## Retailer.countryKorea	3.041567e-08
## Retailer.countryMexico	2.278672e-04
## Retailer.countryNetherlands	2.389600e-04
## Retailer.countrySingapore	1.557950e-05
## Retailer.countrySpain	2.666193e-03
## Retailer.countrySweden	7.483932e-01
## Retailer.countrySwitzerland	9.651452e-01
## Retailer.countryUnited Kingdom	2.261252e-09
## Retailer.countryUnited States	7.577131e-20
## Product.lineGolf Equipment	2.611906e-01
## Product.lineMountaineering Equipment	4.805741e-01
## Product.lineOutdoor Protection	2.589632e-09
## Product.linePersonal Accessories	3.328417e-07
## Year2:Order.method.typeFax	1.494603e-02

## Year2:Order.method.typeMail	2.504252e-01
## Year2:Order.method.typeSales visit	3.792867e-01
## Year2:Order.method.typeSpecial	3.528600e-02
## Year2:Order.method.typeTelephone	1.088002e-01
## Year2:Order.method.typeWeb	4.354632e-09
## Retailer.countryAustria:Product.lineGolf Equipment	7.897408e-01
## Retailer.countryBelgium:Product.lineGolf Equipment	7.559946e-01
## Retailer.countryBrazil:Product.lineGolf Equipment	5.812776e-01
## Retailer.countryCanada:Product.lineGolf Equipment	2.757051e-01
## Retailer.countryChina:Product.lineGolf Equipment	6.278193e-01
## Retailer.countryDenmark:Product.lineGolf Equipment	7.135774e-01
## Retailer.countryFinland:Product.lineGolf Equipment	1.222588e-02
## Retailer.countryFrance:Product.lineGolf Equipment	2.907019e-01
## Retailer.countryGermany:Product.lineGolf Equipment	6.451473e-02
## Retailer.countryItaly:Product.lineGolf Equipment	1.439987e-01
## Retailer.countryJapan:Product.lineGolf Equipment	1.516246e-01
## Retailer.countryKorea:Product.lineGolf Equipment	2.050074e-01
## Retailer.countryMexico:Product.lineGolf Equipment	4.472064e-01
## Retailer.countryNetherlands:Product.lineGolf Equipment	3.045286e-01
## Retailer.countrySingapore:Product.lineGolf Equipment	3.298338e-01
## Retailer.countrySpain:Product.lineGolf Equipment	5.358376e-01
## Retailer.countrySweden:Product.lineGolf Equipment	5.071702e-02
## Retailer.countrySwitzerland:Product.lineGolf Equipment	7.636549e-01
## Retailer.countryUnited Kingdom:Product.lineGolf Equipment	1.266223e-01
## Retailer.countryUnited States:Product.lineGolf Equipment	3.530495e-02
## Retailer.countryAustria:Product.lineMountaineering Equipment	1.114743e-02
## Retailer.countryBelgium:Product.lineMountaineering Equipment	3.186334e-01
## Retailer.countryBrazil:Product.lineMountaineering Equipment	6.413388e-03
## Retailer.countryCanada:Product.lineMountaineering Equipment	3.096672e-04
## Retailer.countryChina:Product.lineMountaineering Equipment	2.568097e-08
## Retailer.countryDenmark:Product.lineMountaineering Equipment	3.193477e-01
## Retailer.countryFinland:Product.lineMountaineering Equipment	1.697643e-03
## Retailer.countryFrance:Product.lineMountaineering Equipment	9.098800e-07
## Retailer.countryGermany:Product.lineMountaineering Equipment	1.197631e-06
## Retailer.countryItaly:Product.lineMountaineering Equipment	9.773645e-03
## Retailer.countryJapan:Product.lineMountaineering Equipment	1.858051e-01
## Retailer.countryKorea:Product.lineMountaineering Equipment	2.267429e-03
## Retailer.countryMexico:Product.lineMountaineering Equipment	1.836614e-03
## Retailer.countryNetherlands:Product.lineMountaineering Equipment	5.504069e-05
## Retailer.countrySingapore:Product.lineMountaineering Equipment	6.540796e-04
## Retailer.countrySpain:Product.lineMountaineering Equipment	2.531470e-01
## Retailer.countrySweden:Product.lineMountaineering Equipment	1.215617e-02
## Retailer.countrySwitzerland:Product.lineMountaineering Equipment	2.504481e-02
## Retailer.countryUnited Kingdom:Product.lineMountaineering Equipment	5.162534e-09
## Retailer.countryUnited States:Product.lineMountaineering Equipment	2.983716e-12
## Retailer.countryAustria:Product.lineOutdoor Protection	1.280911e-01
## Retailer.countryBelgium:Product.lineOutdoor Protection	3.365080e-01
## Retailer.countryBrazil:Product.lineOutdoor Protection	3.020705e-04
## Retailer.countryCanada:Product.lineOutdoor Protection	8.087130e-05
## Retailer.countryChina:Product.lineOutdoor Protection	1.609214e-06
## Retailer.countryDenmark:Product.lineOutdoor Protection	5.584849e-01
## Retailer.countryFinland:Product.lineOutdoor Protection	9.024602e-05
## Retailer.countryFrance:Product.lineOutdoor Protection	2.627046e-06
## Retailer.countryGermany:Product.lineOutdoor Protection	4.561846e-06

## Retailer.countryItaly:Product.lineOutdoor Protection	4.113506e-03
## Retailer.countryJapan:Product.lineOutdoor Protection	5.494774e-08
## Retailer.countryKorea:Product.lineOutdoor Protection	9.999405e-06
## Retailer.countryMexico:Product.lineOutdoor Protection	4.763431e-04
## Retailer.countryNetherlands:Product.lineOutdoor Protection	2.374757e-03
## Retailer.countrySingapore:Product.lineOutdoor Protection	6.202523e-05
## Retailer.countrySpain:Product.lineOutdoor Protection	2.216214e-02
## Retailer.countrySweden:Product.lineOutdoor Protection	3.650906e-01
## Retailer.countrySwitzerland:Product.lineOutdoor Protection	9.889575e-01
## Retailer.countryUnited Kingdom:Product.lineOutdoor Protection	3.332935e-08
## Retailer.countryUnited States:Product.lineOutdoor Protection	1.338604e-15
## Retailer.countryAustria:Product.linePersonal Accessories	2.341201e-03
## Retailer.countryBelgium:Product.linePersonal Accessories	3.035411e-02
## Retailer.countryBrazil:Product.linePersonal Accessories	3.126916e-05
## Retailer.countryCanada:Product.linePersonal Accessories	1.121254e-06
## Retailer.countryChina:Product.linePersonal Accessories	3.844531e-10
## Retailer.countryDenmark:Product.linePersonal Accessories	5.614518e-01
## Retailer.countryFinland:Product.linePersonal Accessories	1.099022e-04
## Retailer.countryFrance:Product.linePersonal Accessories	8.135835e-07
## Retailer.countryGermany:Product.linePersonal Accessories	1.040898e-07
## Retailer.countryItaly:Product.linePersonal Accessories	3.120227e-03
## Retailer.countryJapan:Product.linePersonal Accessories	5.914974e-08
## Retailer.countryKorea:Product.linePersonal Accessories	9.636048e-04
## Retailer.countryMexico:Product.linePersonal Accessories	2.076238e-02
## Retailer.countryNetherlands:Product.linePersonal Accessories	2.003193e-03
## Retailer.countrySingapore:Product.linePersonal Accessories	4.636216e-03
## Retailer.countrySpain:Product.linePersonal Accessories	6.873169e-03
## Retailer.countrySweden:Product.linePersonal Accessories	7.747728e-02
## Retailer.countrySwitzerland:Product.linePersonal Accessories	3.443792e-01
## Retailer.countryUnited Kingdom:Product.linePersonal Accessories	4.584481e-08
## Retailer.countryUnited States:Product.linePersonal Accessories	1.725248e-18

```

# stargazer2(model1, model2, model3)
model3_predictions <- predict(model3, data_200607[, params],
                               interval = "prediction")
cbind(BIC(model1, model2, model3),
      AIC = AIC(model1, model2, model3)[, 2])

```

## df BIC AIC	
## model1 16 310199.3 310078.4	
## model2 40 308373.7 308071.6	
## model3 120 308630.5 307724.2	

```

rbind("model1" = accuracy(model1_predictions[, 1],
                           data_200607$Revenue)[, c("RMSE", "MAE")],
      "model2" = accuracy(model2_predictions[, 1],
                           data_200607$Revenue)[, c("RMSE", "MAE")],
      "model3" = accuracy(model3_predictions[, 1],
                           data_200607$Revenue)[, c("RMSE", "MAE")])

```

## RMSE MAE	
## model1 19712.78 9370.763	
## model2 18265.22 9550.499	
## model3 17994.27 9672.254	

The BIC is larger in the 3rd model (including the interaction terms) because the reduction in the errors is not compensated by the extra complexity introduced by adding the interaction terms. Regarding to the forecasts, the RMSE has kept decreasing with each new model, but the MAE has increased. This means that the prediction errors are higher on average, but their variance is smaller (the RMSE penalizes observations that differ largely from the real value).

- Is the change in the average revenue different from 95 cents when the planned revenue increases by \$1?

```
params = c("Planned.revenue")
model4 <- lm(as.formula(paste("Revenue", paste(params, sep = "", collapse = " + ")),
                         sep = " ~ ")), data_200405)
coeftest(model4, vcov = vcovHC) [, 1]

##      (Intercept) Planned.revenue
## -3276.6490936     0.9693842

new_data <- data.frame(data_200607[, params])
names(new_data) <- params
```

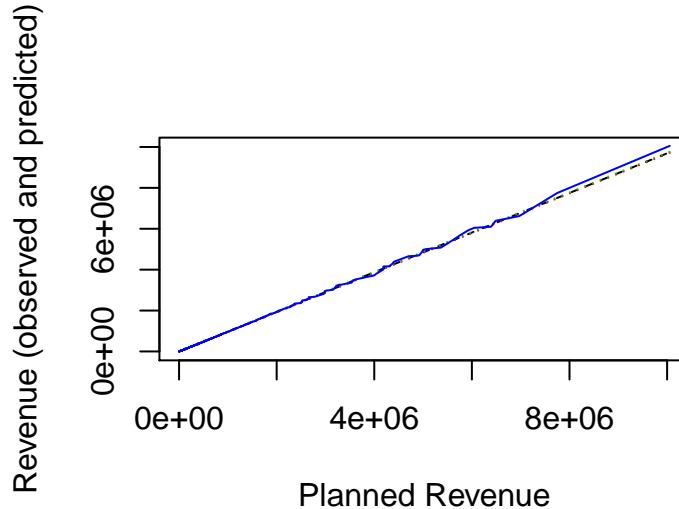


Figure 27: Planned Revenue vs. Revenue (observed and predicted) in 2006 and 2007

```
BIC(model4); AIC(model4)

## [1] 310236.7
## [1] 310214
```

```
(RMSE <- sqrt(mean((model4_predictions[, 1] - data_200607$Revenue)^2)))  
  
## [1] 19593.67  
  
accuracy(model4_predictions[, 1], data_200607$Revenue) [, c("RMSE", "MAE")]  
  
##      RMSE      MAE  
## 19593.667 8804.353
```

Based on its RMSE and MAE, this model is similar or even better than the ones we have built so far.

```
linearHypothesis(model4, "Planned.revenue = 0.95", vcov = vcovHC)  
  
## Linear hypothesis test  
##  
## Hypothesis:  
## Planned.revenue = 0.95  
##  
## Model 1: restricted model  
## Model 2: Revenue ~ Planned.revenue  
##  
## Note: Coefficient covariance matrix supplied.  
##  
##   Res.Df Df      F    Pr(>F)  
## 1 14086  
## 2 14085  1 50.771 1.088e-12 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
  
linearHypothesis(model4, paste("Planned.revenue =",  
                               coefest(model4, vcov = vcovHC)[2, 1]),  
                               vcov = vcovHC)  
  
## Linear hypothesis test  
##  
## Hypothesis:  
## Planned.revenue = 0.969384229459145  
##  
## Model 1: restricted model  
## Model 2: Revenue ~ Planned.revenue  
##  
## Note: Coefficient covariance matrix supplied.  
##  
##   Res.Df Df  F Pr(>F)  
## 1 14086  
## 2 14085  1  0      1
```

As shown below, the change in the average revenue is significantly different from \$0.95 when the revenue increases by \$1 (while the  $F$  statistic of the exact value, which is quite close to \$0.95, has a  $p$  value equal to 1):

- Explain what interaction terms in your model mean in context supported by data visualizations.

For example, in the 1st model (where the baseline is 2004 and orders arrived by e-mail), an additional year increases the revenues, on average, by about \$2,871, and an order that arrived by fax is, also on average, about \$4,038 higher. But if both conditions happen at the same time (i.e., an order by fax in 2005), the effect of the subsequent year disappears (because the interaction term estimate is  $-\$2,875$ ). Similarly, in the 2nd model, where the baseline is Australia and Camping Equipment, the revenue from Germany is, on average, about \$2,513 lower, and one for Golf Equipment is \$5,623 lower. Putting both conditions together would decrease revenue by \$8,136, but when the interaction term is included (in the 3rd model), the mentioned estimates change to  $-\$4,472$  and  $-\$1,925$ , respectively, while the interaction term is  $-\$6,077$ ; i.e., the decrease in revenue is much larger, almost by \$12,500.

- Give two reasons why the OLS model coefficients may be biased and/or not consistent, be specific.
- Propose (but do not actually implement) a plan for an IV approach to improve your forecasting model.

Let's say we are able to know in advance the quantity of each product that will be purchased in the coming years, and thus incorporate that variable in our model and make predictions based on that new model. Another variable that would certainly help forecast the revenue is the **demand** (i.e., not what is actually sold but what customers are willing to buy, but may end up not doing so for a number of reasons: they prefer other alternatives, they don't have the money, they don't even know our product, etc.). Demand cannot be included as an explanatory variable (i.e., it's omitted from our model), so it's included in the error term. Since the actual quantity that is bought is correlated with demand, our estimates may be biased. A possible IV (that has no effect on **revenue**—after controlling for **quantity**—, is not correlated with demand, and is correlated with **quantity**) could be a metric of the **brand awareness**. Or an indicator of the economic situation.

---